

EECS 545: Machine Learning

Lecture 1. Introduction

Honglak Lee

1/10/2020



Outline

- Administrative
- What is machine learning?

Teaching staff

- Instructor: Honglak Lee
 - Email: honglak@eecs.umich.edu
 - Office: CSE 3773
- Graduate Student Instructor:
 - Sungryull Sohn (srsohn@umich.edu)
 - Yijie Guo (guoyijie@umich.edu)
 - GS Oh (gsoh@umich.edu)
- For office hours, please see the future announcement and calendar.
- For all questions, please use Piazza:
 - <https://piazza.com/class/jqn72gu5i3v4l4>

About this course

- Graduate-level introduction of machine learning
- Foundations of machine learning
 - Mathematical derivation, Implementation of the algorithms, applications
- Topics
 - supervised learning
 - unsupervised learning
 - reinforcement learning
- Other topics
 - deep learning, learning theory, probabilistic models, sparsity and feature selection, Bayesian techniques, ensemble methods

About this course

- Cover practical applications of machine learning
 - computer vision, data mining, speech recognition, text processing, robot perception and control, etc.
- Our goal is to help you to
 - Understand fundamentals of machine learning
 - Learn technical details of ML algorithms
 - Learn how to implement some important algorithms
 - Use machine learning algorithms for your research and applications.

Text books

There will be no official textbook for the course. However, the following materials will be helpful:

- Chris Bishop, "Pattern Recognition and Machine Learning". Springer, 2007.
- Kevin Murphy, "Machine Learning: A Probabilistic Perspective", 2012.
- David Barber, "Bayesian Reasoning and Machine Learning", 2017. (available online) <http://web4.cs.ucl.ac.uk/staff/D.Barber/pmwiki/pmwiki.php?n=Brml.Online>
- Ian Goodfellow and Yoshua Bengio and Aaron Courville, "Deep Learning", (available online) <http://www.deeplearningbook.org/>
- Hastie, Tibshirani, Friedman, "Elements of Statistical Learning". Springer, 2010. (available online) <http://www-stat.stanford.edu/~tibs/ElemStatLearn/>
- Sutton and Barto, "Reinforcement Learning: An Introduction," MIT Press, 2018 (available online) <http://incompleteideas.net/book/the-book-2nd.html>
- (optional) Boyd and Vandenberghe, "Convex Optimization," Cambridge University Press, 2004. (available online) <http://www.stanford.edu/~boyd/cvxbook/>
- (optional) Mackay, "Information Theory, Inference, and Learning Algorithms". Cambridge University Press. 2003. (available online) <http://www.inference.phy.cam.ac.uk/itprnn/book.pdf>

Prerequisites

- Undergrad linear algebra (e.g., MATH 217, MATH 417)
- Multivariate calculus
- Undergrad probability and statistics (e.g., EECS 401)
- Programming skills (equivalent to EECS 280, EECS 281, and experience in python)
 - Nontrivial level of programming is required.
- EECS 492: Introduction to AI
 - (advisory: you can still take this course without completing 492)
- NOTE: If you **have not** taken **at least two** of linear algebra, multivariate calculus, and probability courses, it is **strongly recommended** that you finish them first before taking this course.

Grading policy

- Homework: 30%
- Midterm: 30% (tentative date: late March/early April)
 - [Honor code](#)
- Project: 40%
 - progress report (10%)
 - final report (30%)
- Note: There will be no final exam.
- Extra credits: Up to 2% may be awarded for participation (in class and piazza).

Language of Choice: **Python**

- **Python** is a great language overall for machine learning, with modern libraries and excellent online tutorials for various tasks
- We will utilize **Python** throughout the course, and we encourage you to do the same.
 - We will be using popular libraries, such as **numpy** and **pytorch**.
- There will be a tutorial session (date: TBD)

Homework

- There will be 5 or 6 problem sets.
- Goal: strengthen the understanding of the fundamental concepts, mathematical formulations, algorithms, detailed implementations, and the applications.
- The problem sets will also include programming assignments to implement algorithms covered in the class.
- Homework #1 will be out next Tue (1/14) – due 1/28, 5pm via Gradescope (linked to Canvas).

Late days

- 3 maximum late days per assignment
 - No homework will be accepted 3 days after the due date.
- Total 8 late days allowed.
- After using up all late days, your assignment will be penalized by 20% from your scores.

Study group

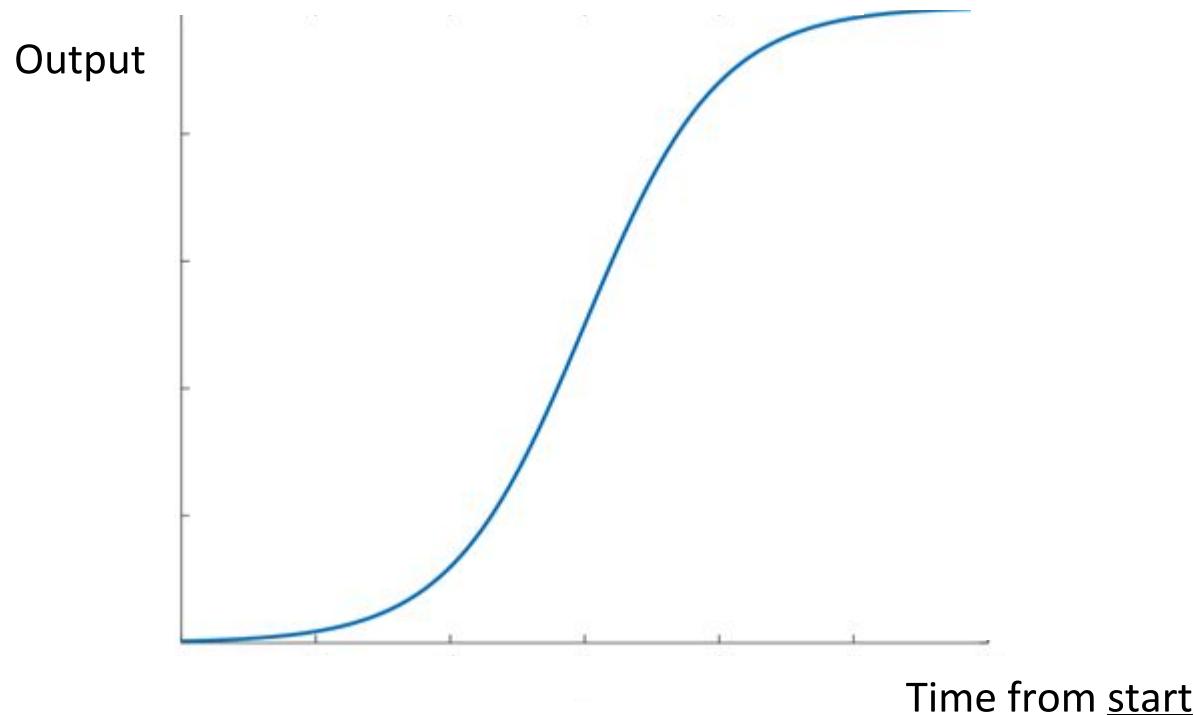
- Form your study group early on!
 - Up to four people are allowed.
- For homework, you may discuss between the study group members, but you should write your own solution independently.
- In the homework submissions, you must put:
 - the names of other people you collaborated.
 - submission time.
- Please start on homework early. (Warning: cramming does not work!)

Course Project

- Scope
 - develop new algorithms and theory in machine learning
 - apply existing algorithms to new problems
 - apply to your own problems of interest
- Milestones (tentative)
 - project proposals (tentative due date: 2/4)
 - progress reports (tentative due date: 3/10)
 - poster presentations (tentative date: 4/24)
 - Final project report (tentative due date: 4/27)
- Requirement
 - Write a 8-page paper
 - Submit the final code
 - Give a poster presentation
- Evaluation is based on:
 - novelty, technical quality, significance, and presentation quality of the project.

Course Project

- 4 or 5 people can form a project group.
- Talk to instructor if you want to get suggestions about project topics.
- Start early!!! (form your group and start working)



Other Information

- Review sessions
 - Will hold review sessions on background materials (linear algebra, probability, Python.)
- Beginning-of-course survey: <https://bit.ly/39XZ5ZY>

Any questions?

Outline

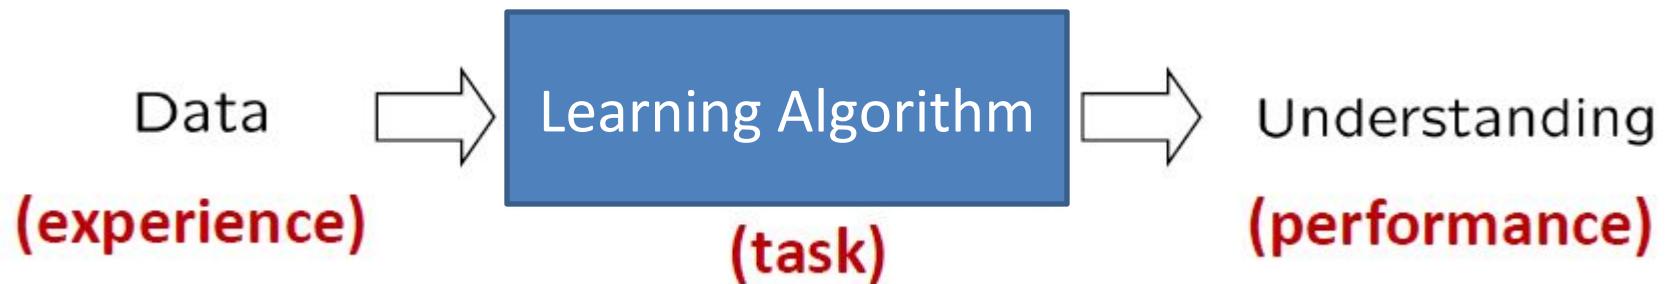
- Administrative
- What is machine learning?

Definition of Machine Learning

- Formal definition (Mitchell 1997): A computer program **A** is said to **learn from experience E** with respect to some class of **tasks T** and **performance measure P**, if its performance at tasks in **T**, as measured by **P**, improves with experience **E**.

Informal definition

- Algorithms that improve their prediction performance at some task with experience (or data).



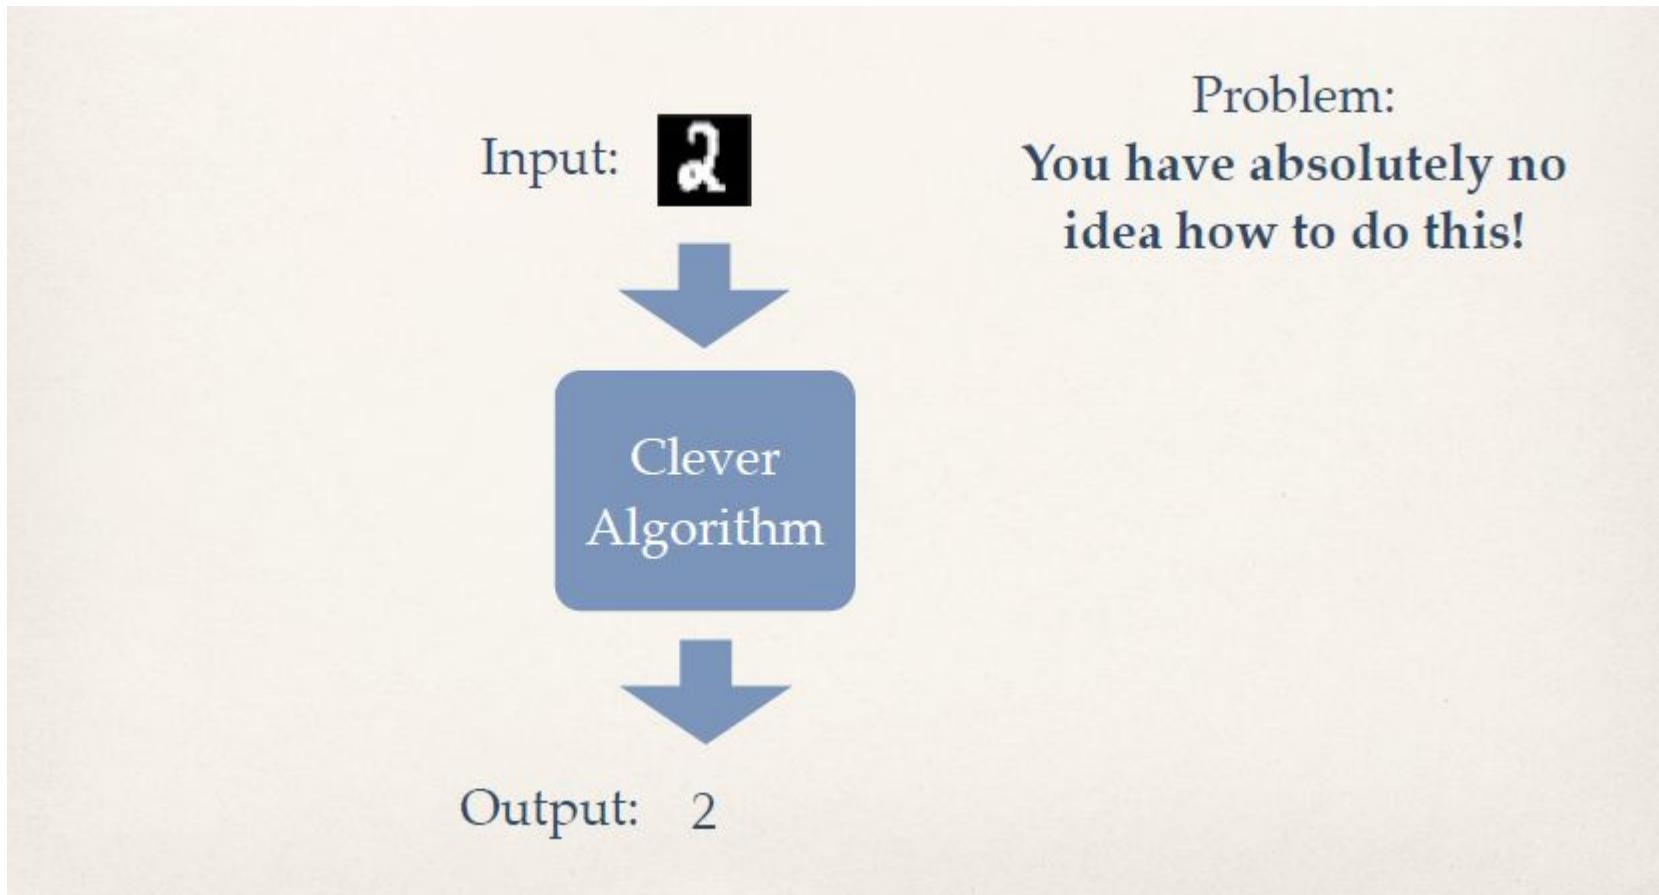
Example: Spam email filtering

“A computer program is said to *learn* from experience E with respect to some task T and some performance measure P, if its performance on T, as measured by P, improves with experience E.”

- Suppose your email program watches which emails you do or do not mark as spam, and based on that learns how to better filter spam.
- Task:
 - Classifying emails as spam or not spam.
- Experience:
 - Watching you label emails as spam or not spam.
- Performance measure
 - The number (or fraction) of emails (disjoint from the training emails) correctly classified as spam/not spam.

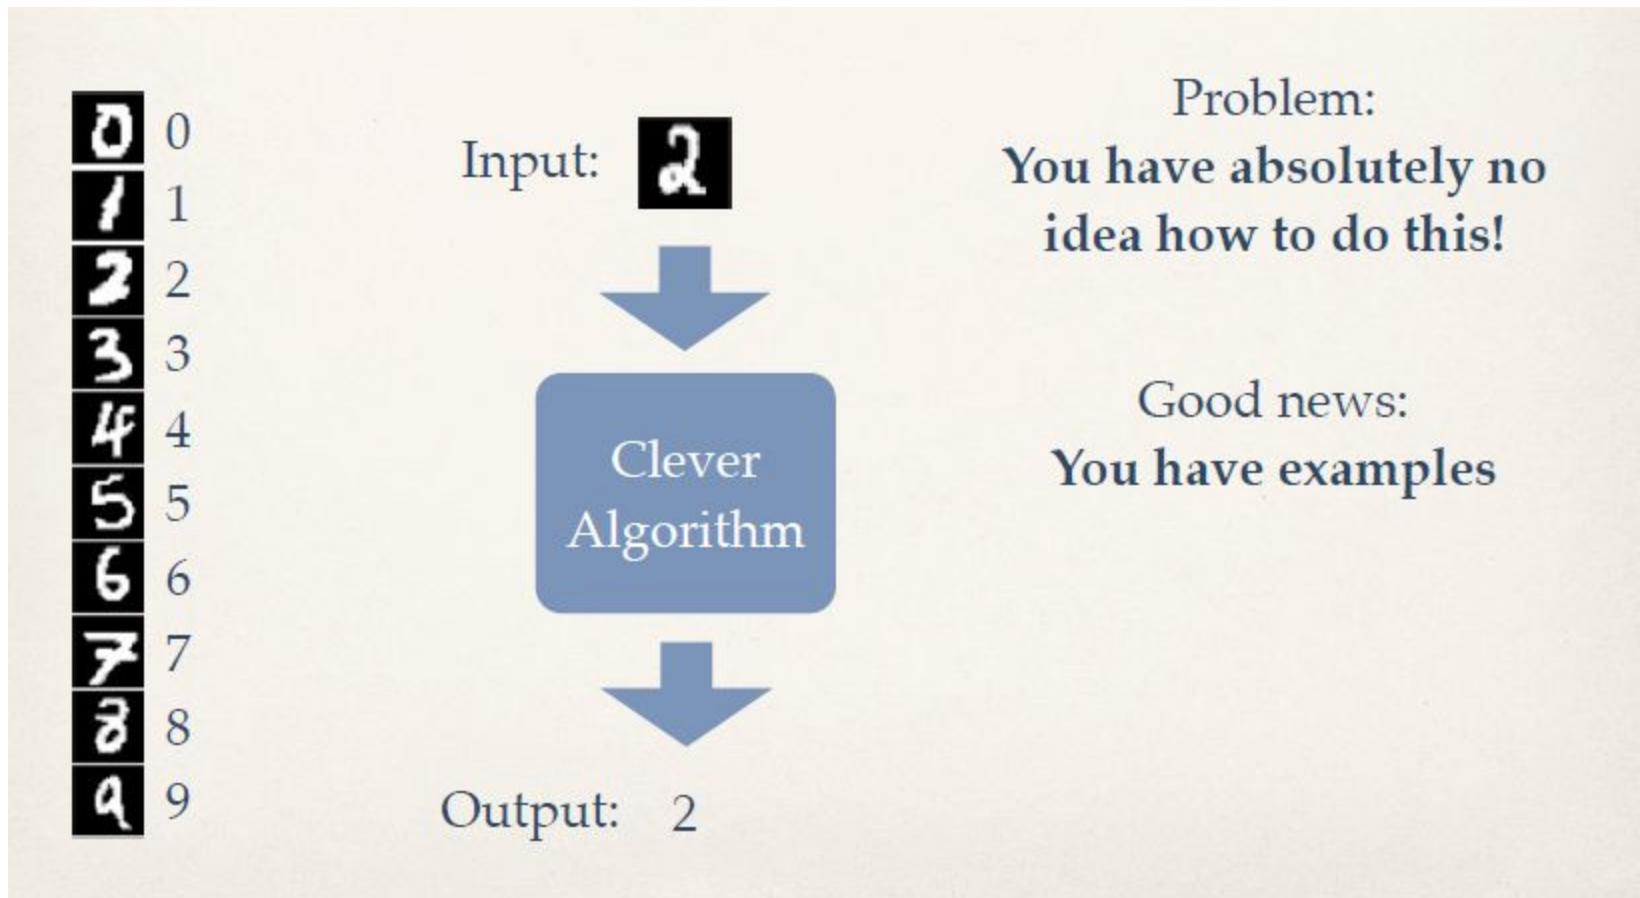
Example

- Problem: Given an image of a handwritten digit, what digit is it?



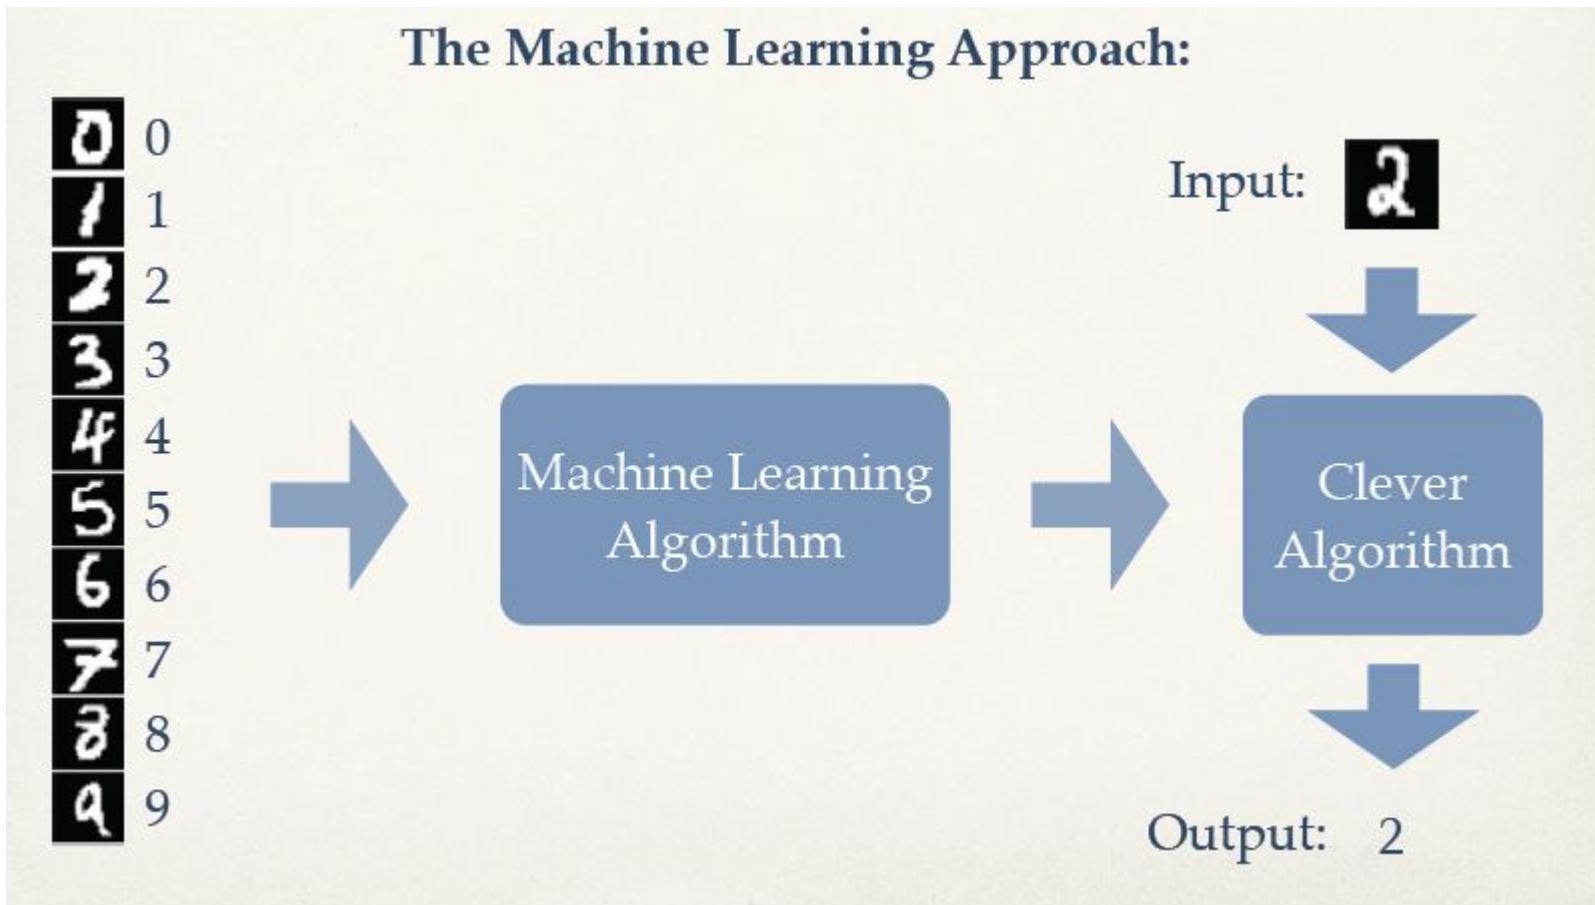
Example

- Problem: Given an image of a handwritten digit, what digit is it?



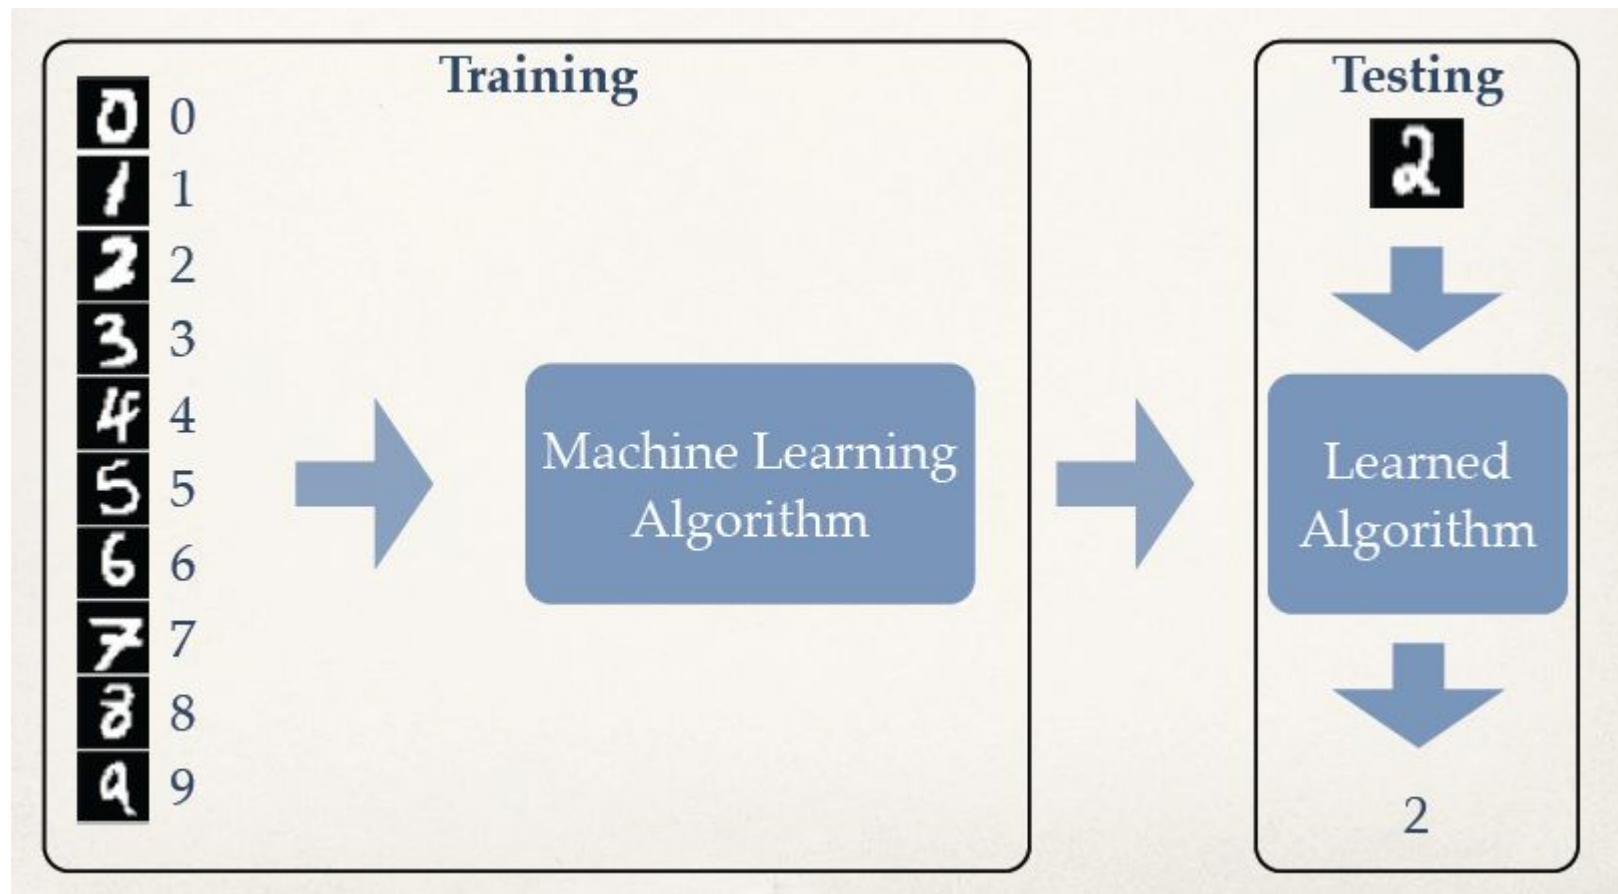
Example

- Problem: Given an image of a handwritten digit, what digit is it?



Example

- Problem: Given an image of a handwritten digit, what digit is it?



Machine Learning Tasks

- Supervised Learning
 - Classification
 - Regression
- Unsupervised Learning
 - Clustering
 - Density estimation
 - Embedding / Dimensionality reduction
- Reinforcement Learning
 - Learning to act (e.g., robot control, decision making, etc.)

Supervised Learning

Given a dataset $D = \{(x_1, y_1), \dots, (x_n, y_n)\}$,
where

- $x_i \in \mathcal{X}$: input (feature)
- $y_i \in \mathcal{Y}$: output (label)

a black box ML algorithm produces a
prediction function $h : \mathcal{X} \rightarrow \mathcal{Y}$, such that
 $h(x)$ can predict the y values for all x
(including training data $x_i \in D$ and unseen
test data x^*).

Supervised Learning

- Labels could be discrete or continuous
 - Discrete labels: **classification**
 - Continuous labels: **regression**

Supervised Learning - Classification

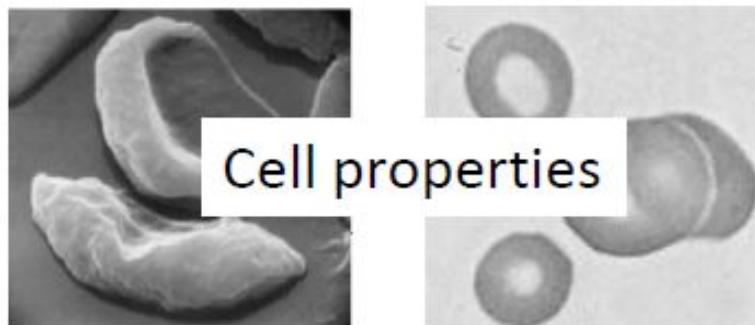
Feature Space \mathcal{X}



Words in a document

Label Space \mathcal{Y}

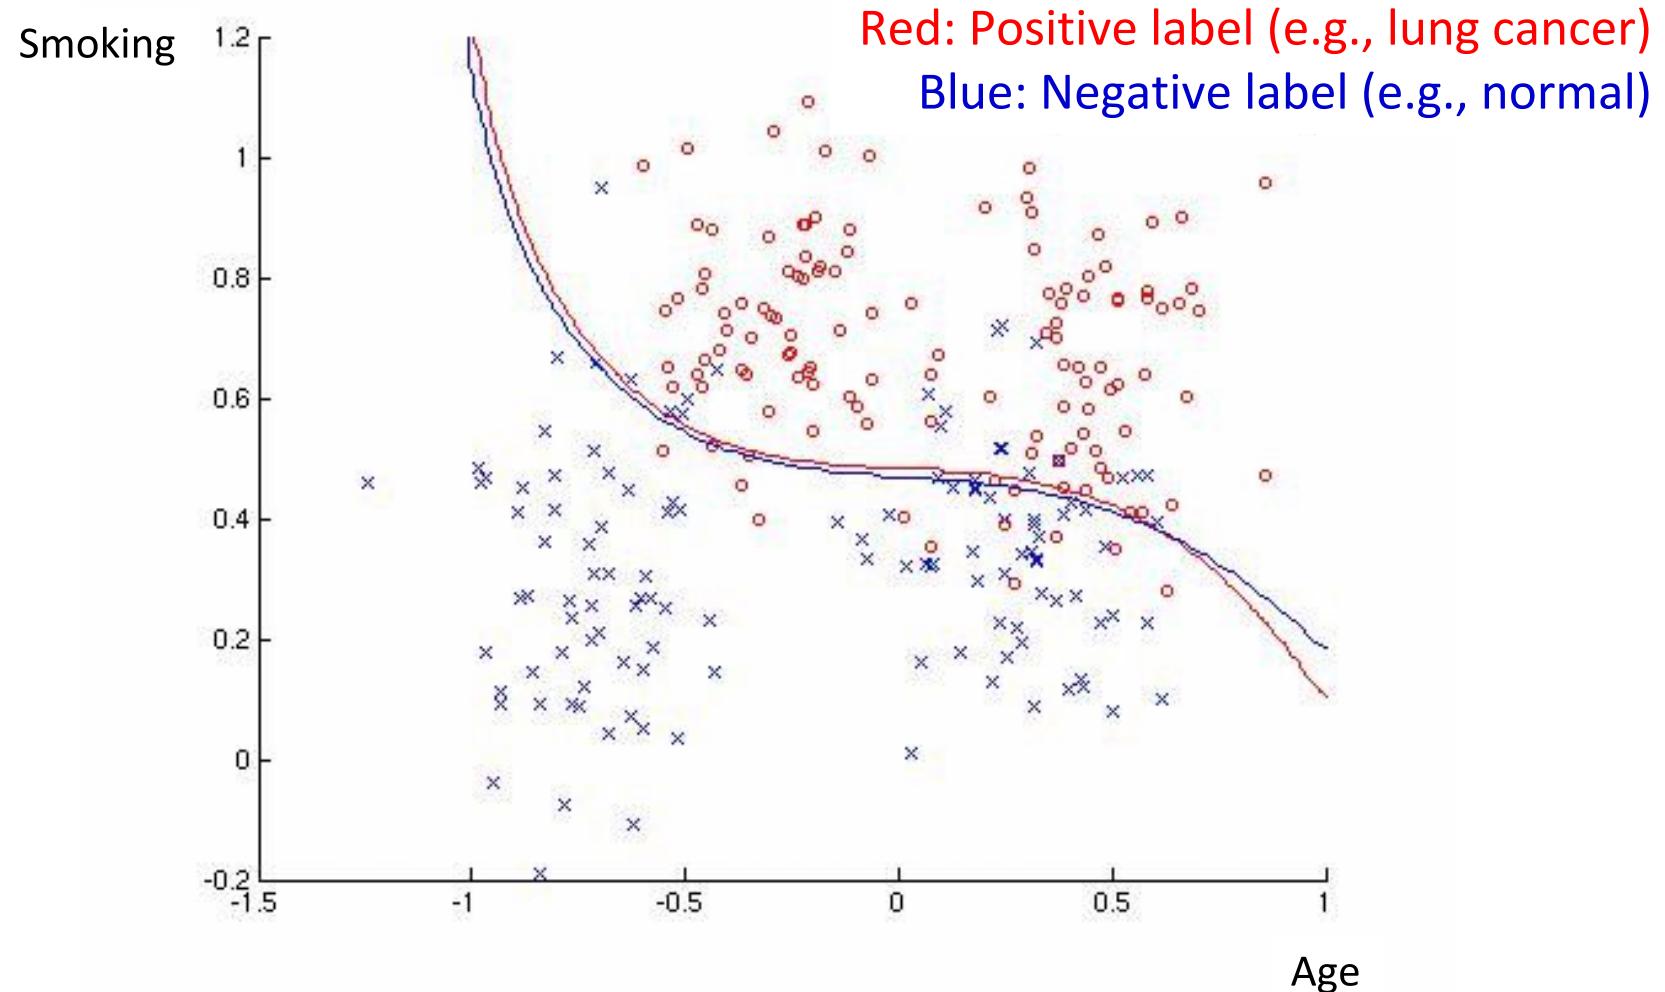
“Sports”
“News”
“Science”
...



“Anemic cell”
“Healthy cell”

Discrete Labels

Supervised Learning - Classification



“Learning decision boundaries”

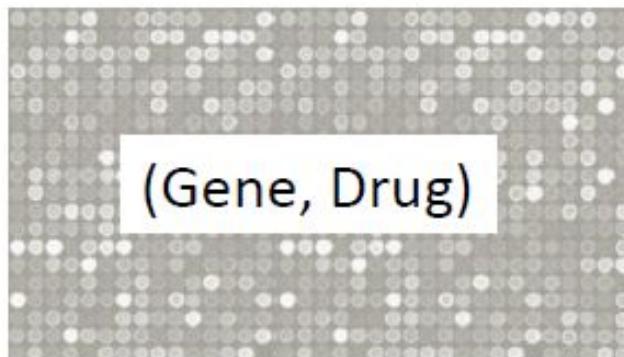
Supervised Learning - Regression

Feature Space \mathcal{X}



Label Space \mathcal{Y}

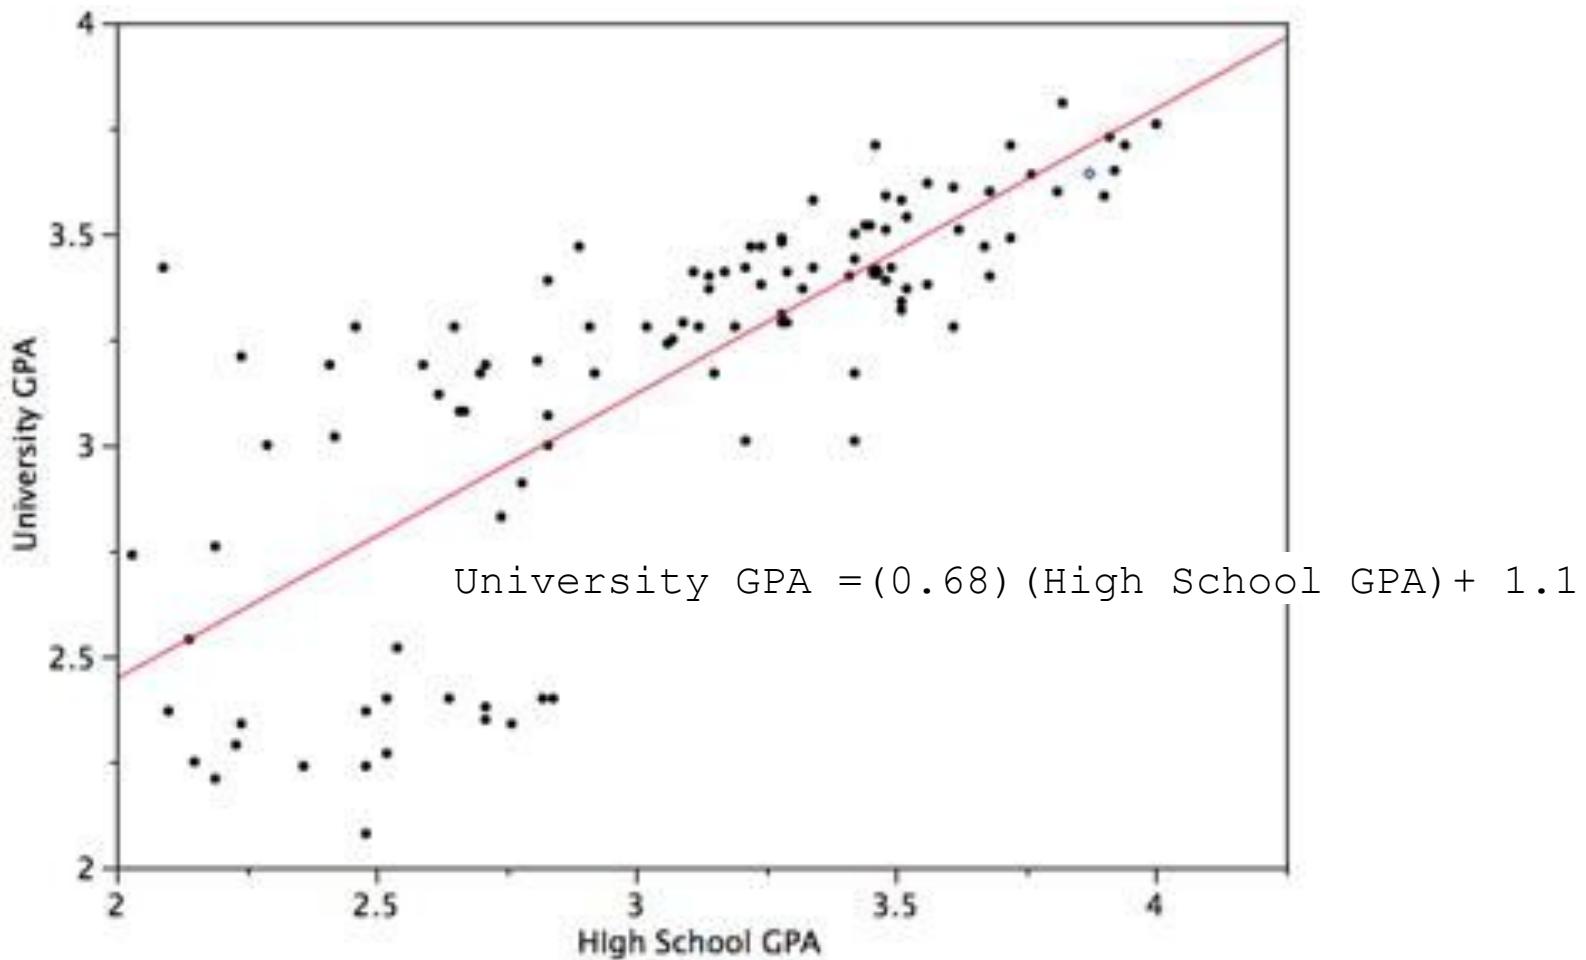
Share Price
“\$ 24.50”



Expression level
“0.01”

Continuous Labels

Supervised Learning - Regression



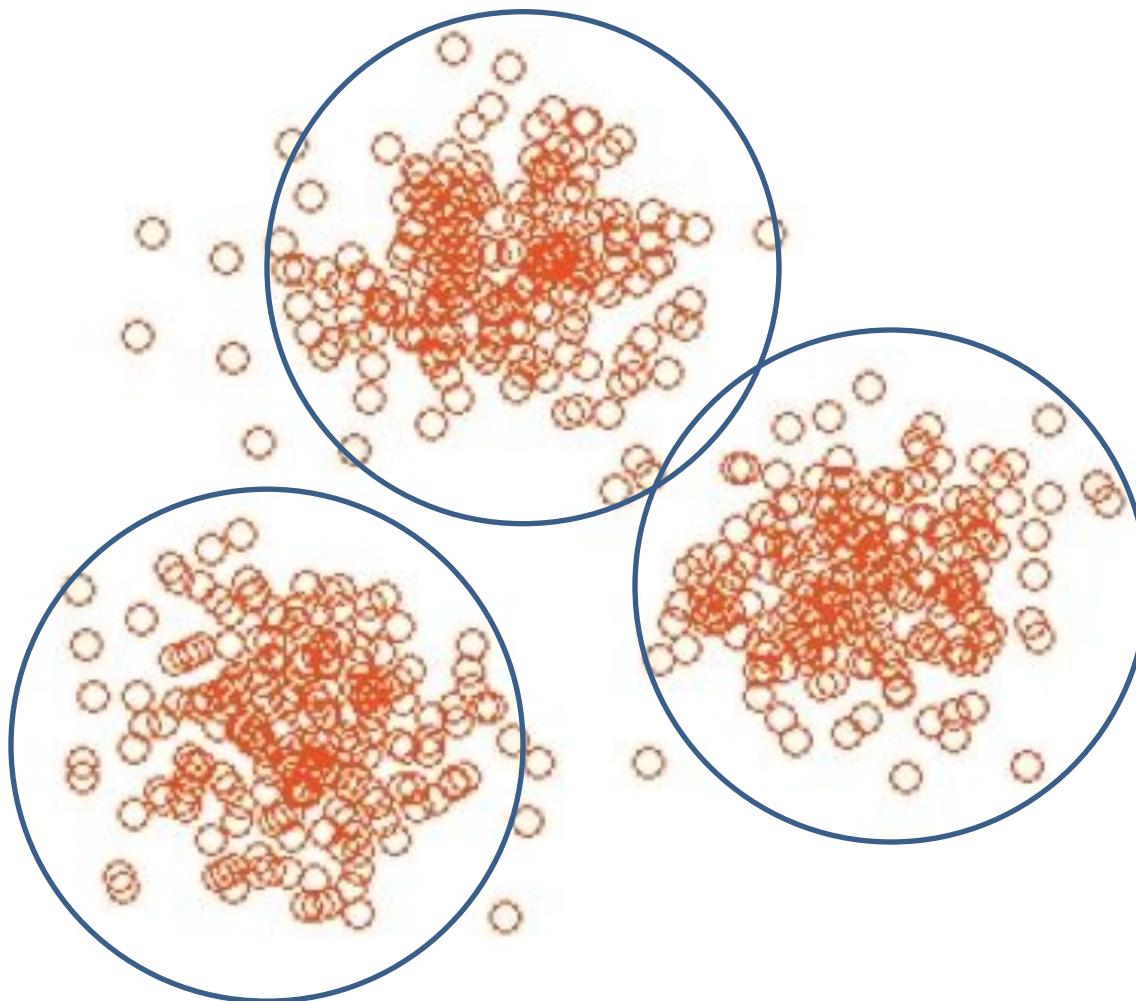
“Learning regression function $f(X)$ ”

Unsupervised Learning

- Goal:
 - Given data X without any labels
 - Learn the **structures** or **distribution** of the data
 - Clustering
 - Probability distribution (density)
 - Generating data
 - Embedding & neighborhood relations
- “Learning without teacher (supervision)”

Unsupervised Learning – Clustering

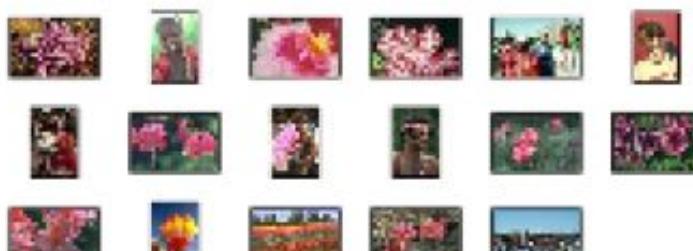
- “Grouping into similar examples”



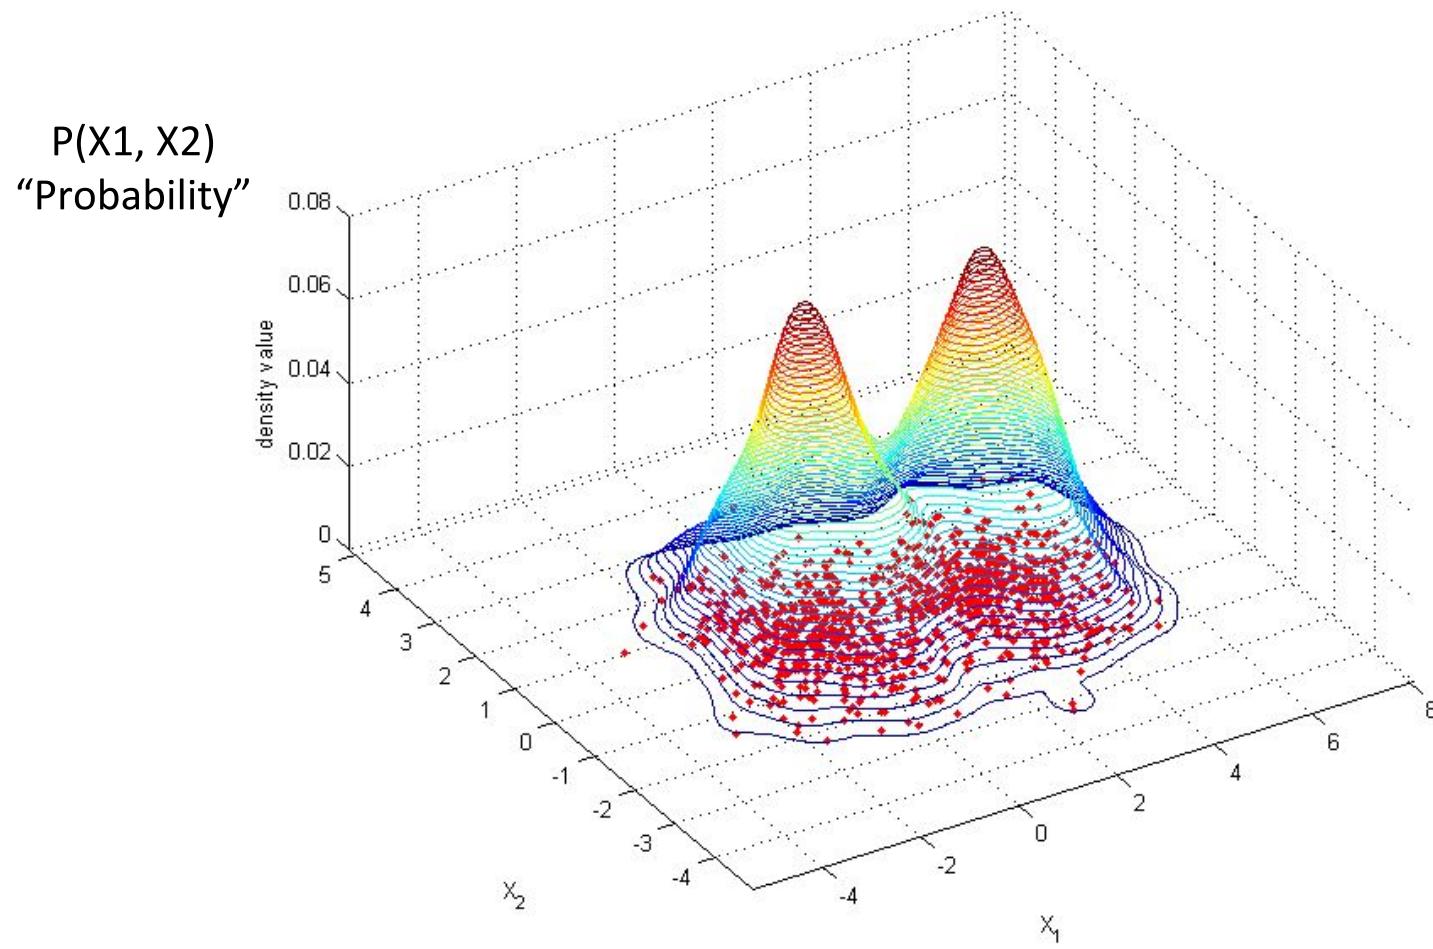
Unsupervised Learning – Clustering

Group similar things e.g. images

[Goldberger et al.]

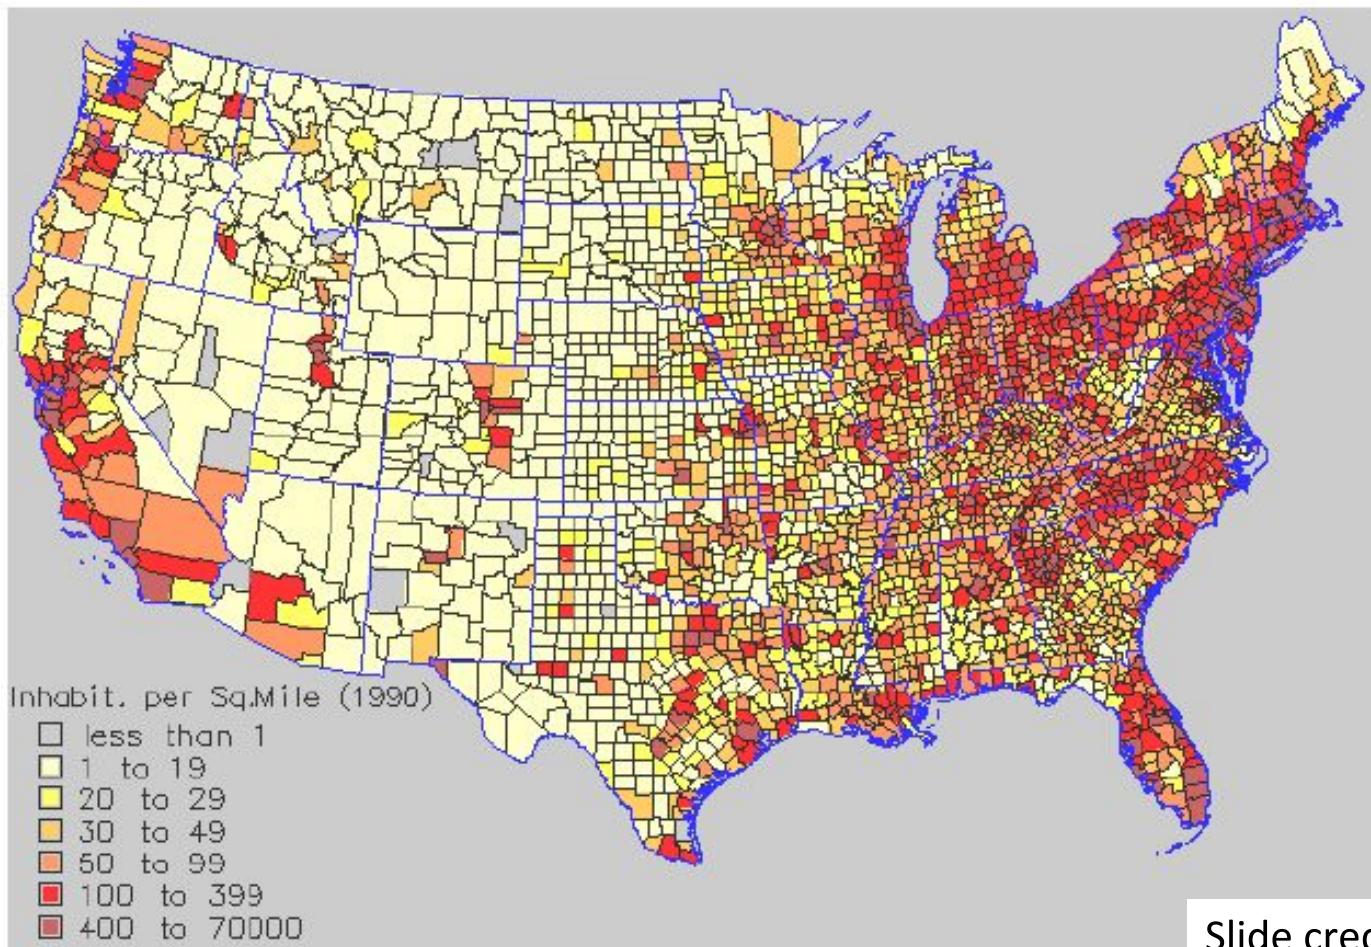


Unsupervised Learning – Density estimation



Unsupervised Learning – Density estimation

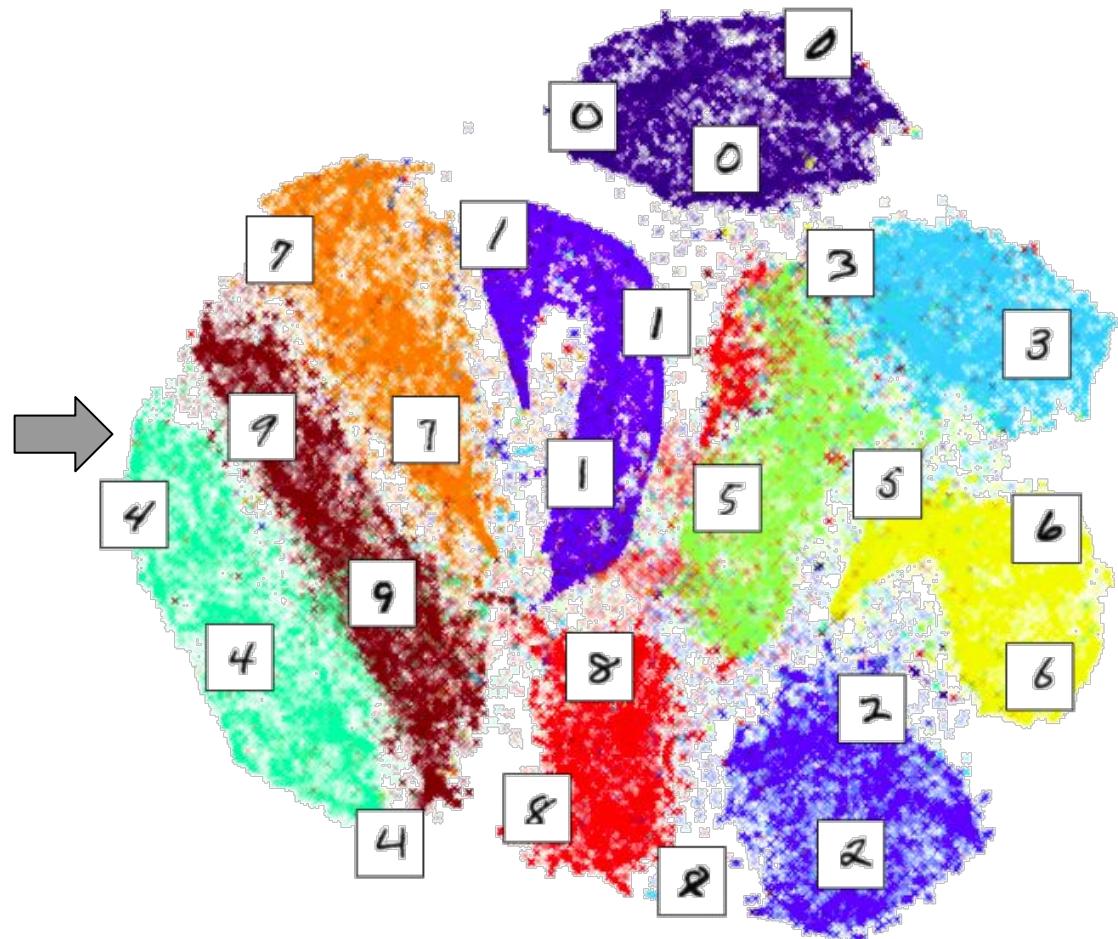
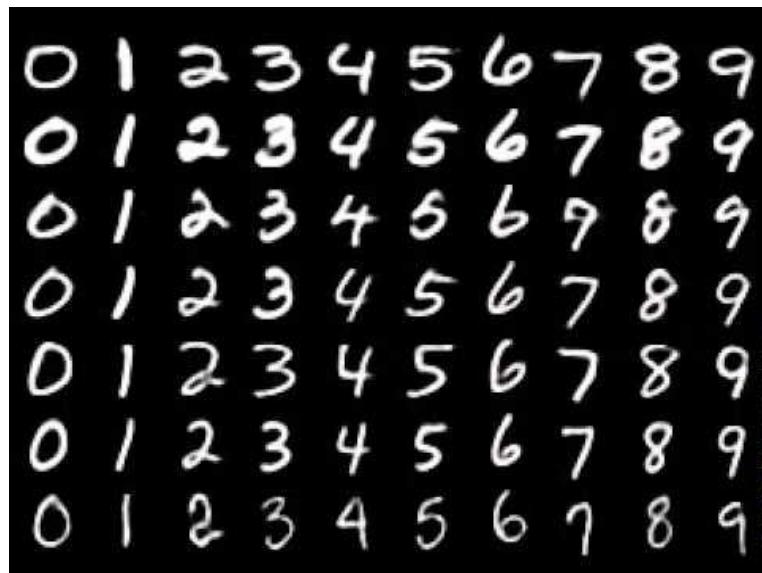
Population density



Slide credit: Aarti Singh

Unsupervised Learning-Embedding and Dimensionality reduction

- E.g., Reducing handwritten digits (784 dim) into low dimensional coordinates



[Maaten and Hinton, 08]

Reinforcement Learning

- Setting
 - Given sequence of states X and “rewards” (e.g., delayed labels)
 - Agent has to take actions A for each time step
- Goal:
 - How to “learn to act” or “make decisions” to maximize the sum of future rewards
- Example: Robot navigation task
 - Input: Dynamical environment + sensor input
 - Action: control signals
 - Rewards: time to reach goal without colliding with obstacles

Reinforcement Learning – learning to control

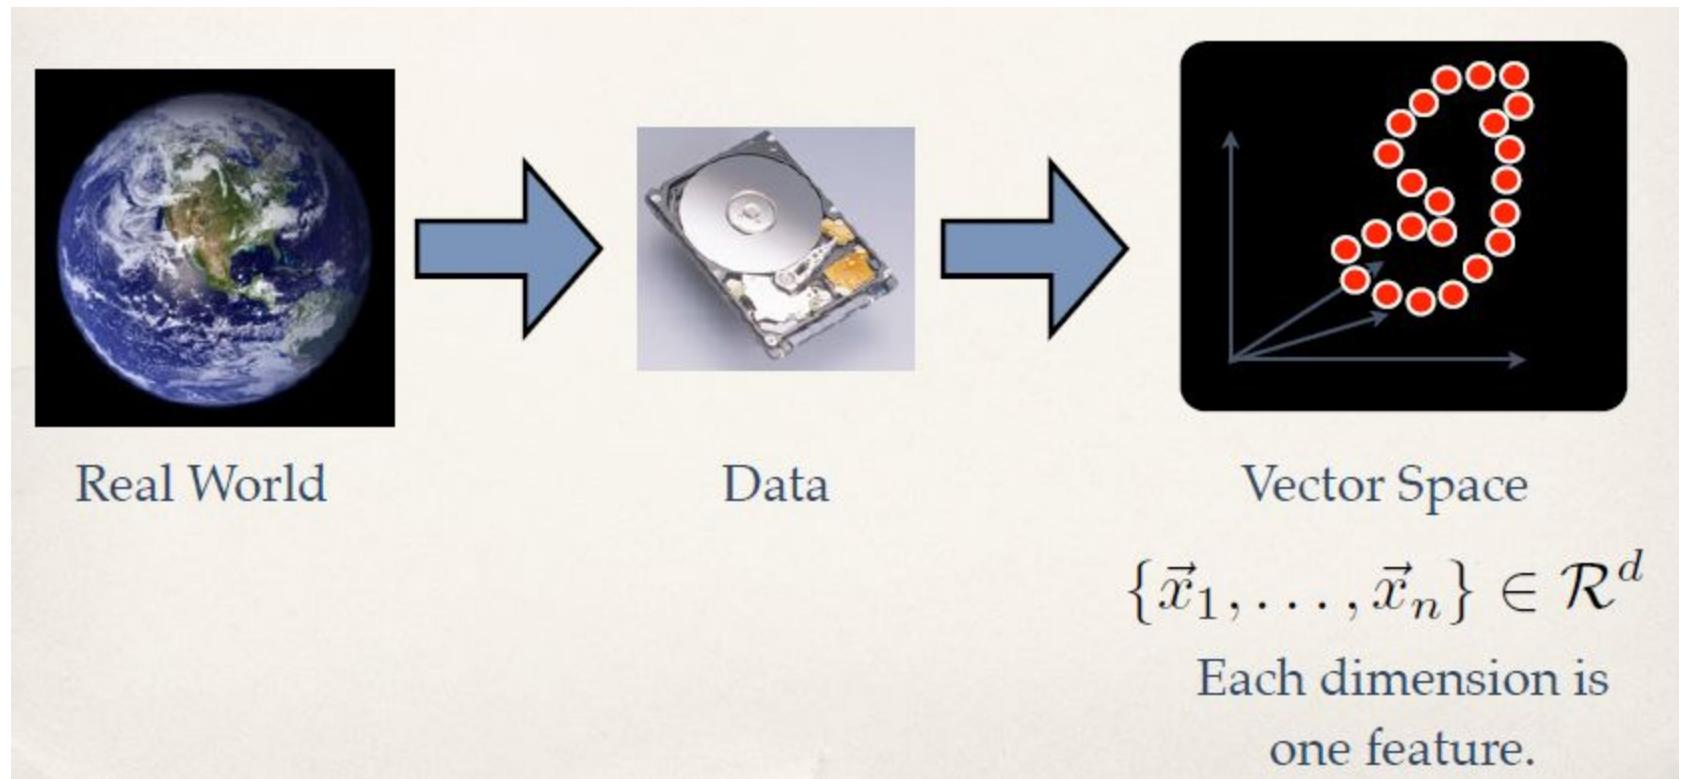
- Example: Robot walking
 - States: sensor inputs, joint angles
 - Action: servo commands for joints
 - Rewards:
 - 1 for reaching the goal
 - -1 for falling down
 - 0 otherwise
- Goal: How can we provide control inputs to maximize the expected future rewards?



Feature representations

Feature Extraction

- Represent data in terms of vectors.
 - Features are **statistics** or **attributes** that describe the data.



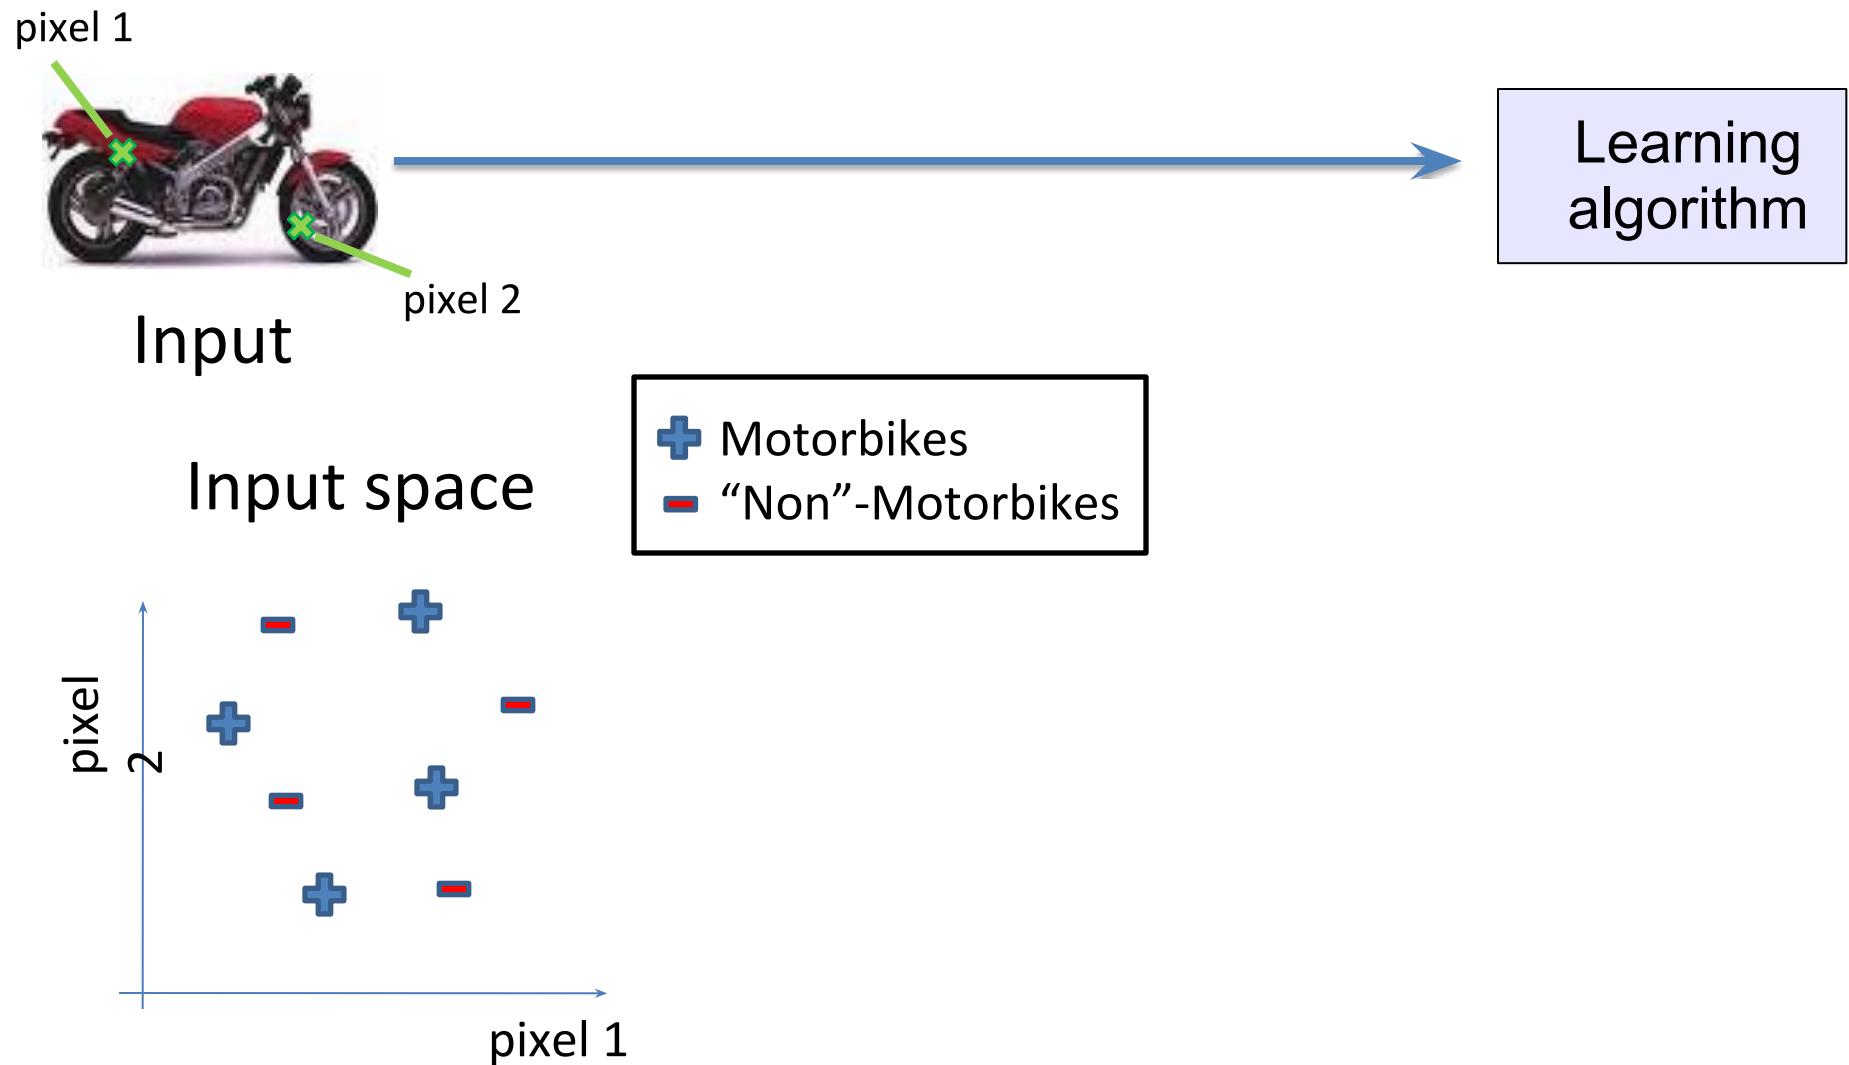
Examples of features: Housing data

- **Given statistics about houses in a local area, predict median value of homes.**
 - #ROOM: average number of rooms per dwelling
 - AREA: average area of house in square foot
 - AGE: proportion of owner-occupied units built prior to 1940
 - CRIME: per capita crime rate by town
 - RESZONE: proportion of residential land zoned for lots over 25,000 sq.ft.
 - INDUS: proportion of non-retail business acres per town
 - NOX: nitric oxides concentration (parts per 10 million)
 -
- **Label: Median value of owner-occupied homes**

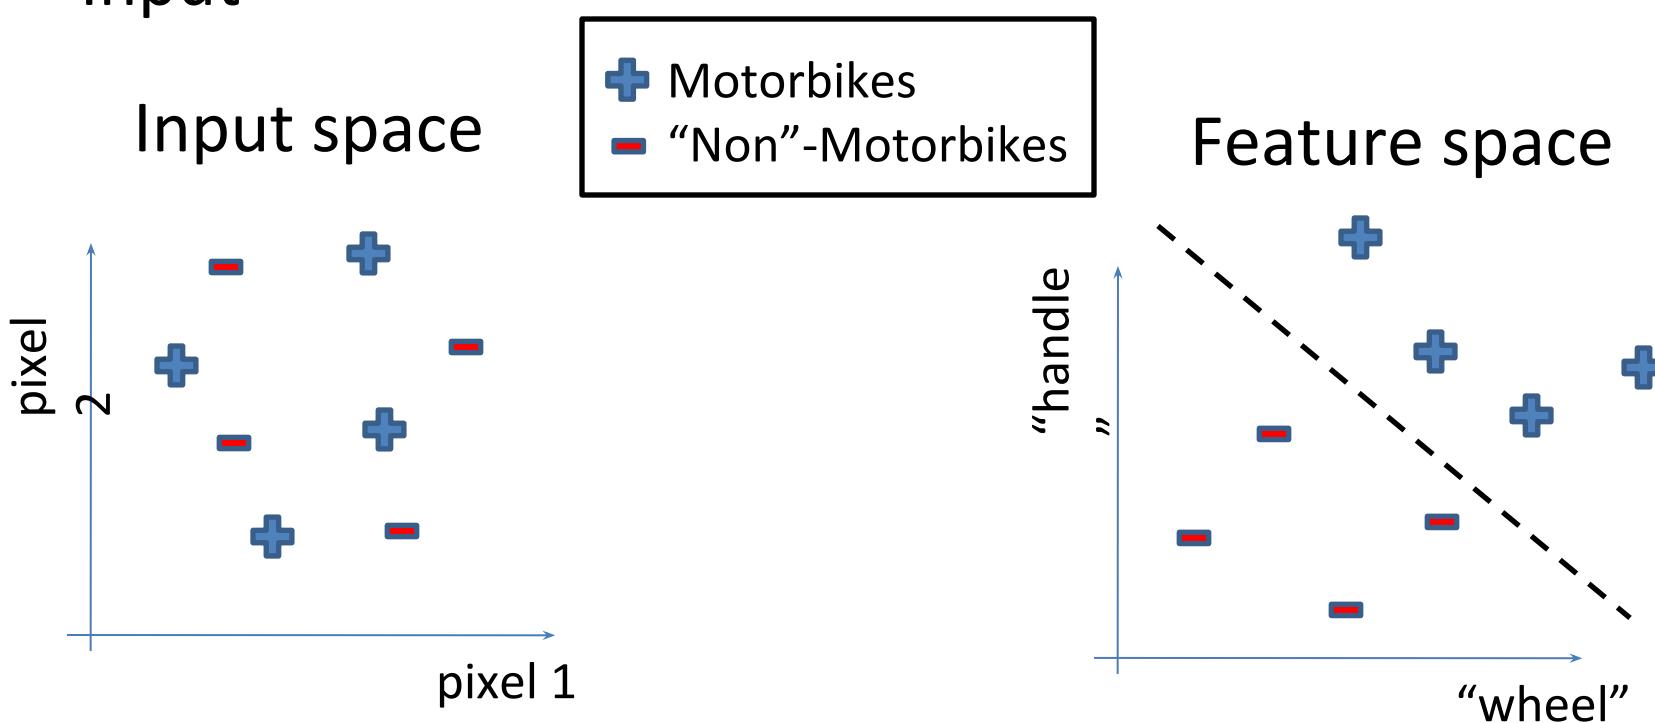
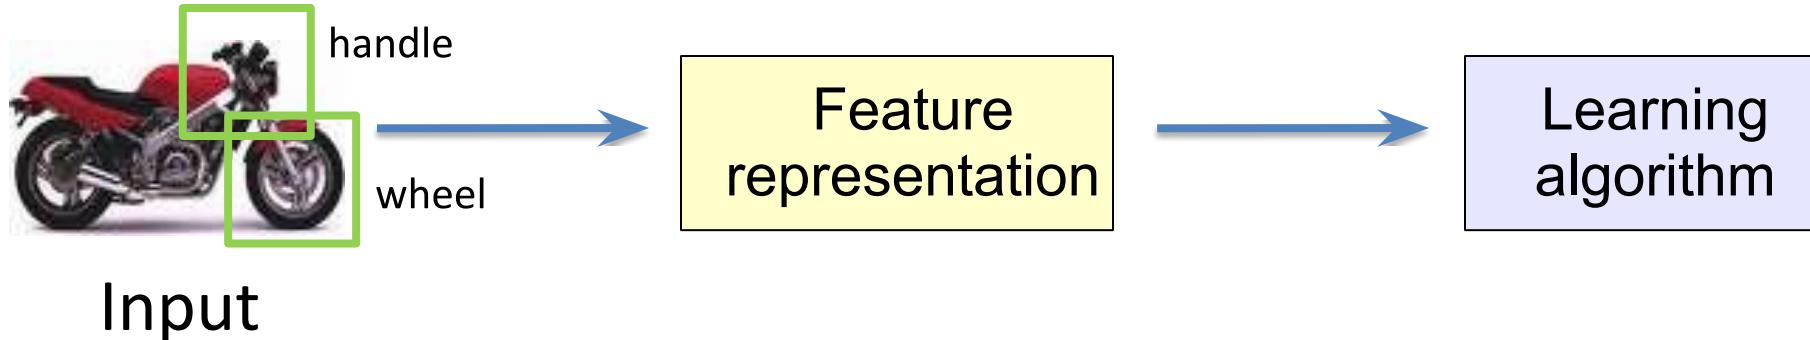
Examples of features: Recognizing handwritten-digits

- Input: 28x28 pixel digit images
- Output: class labels $\in \{0, \dots, 9\}$
- The following basic features can be used:
 - Pixel values (784 dimensional vectors)
 - Aspect ratio of the tight bounding boxes
 - Existence of long vertical strokes
 - Existence of long horizontal strokes
 -

Learning pipeline



Learning pipeline



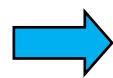
(Traditional) Computer Perception Pipeline



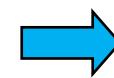
Object
detection



Imag
e

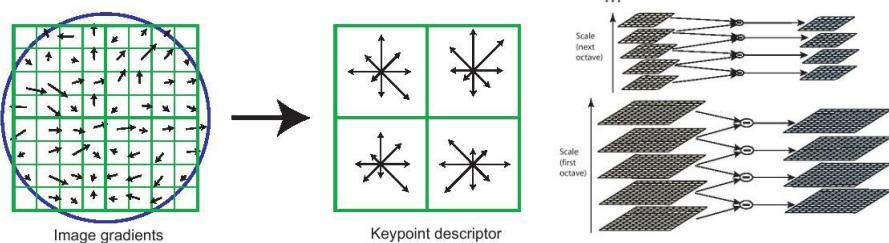


Low-level
vision
features

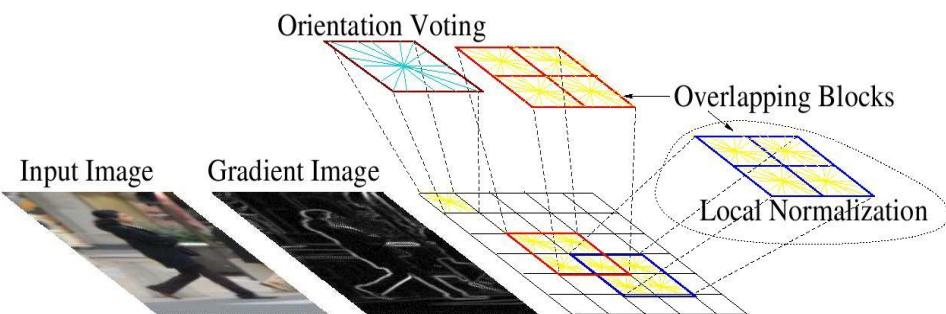


Recognitio
n

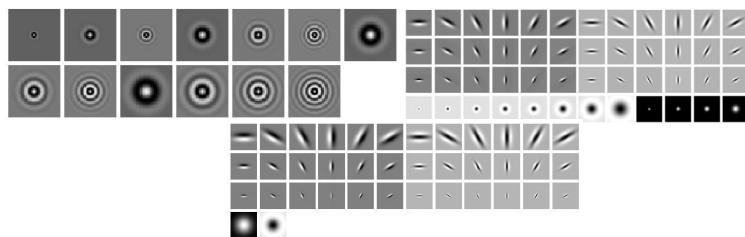
(Traditional) Computer vision features



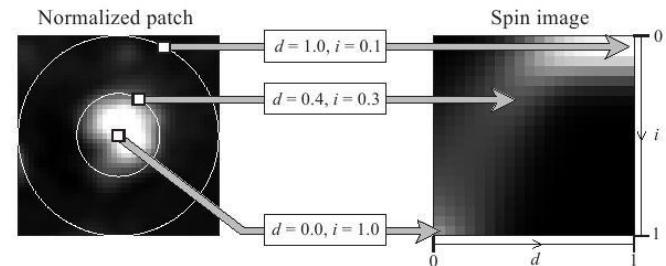
SIFT



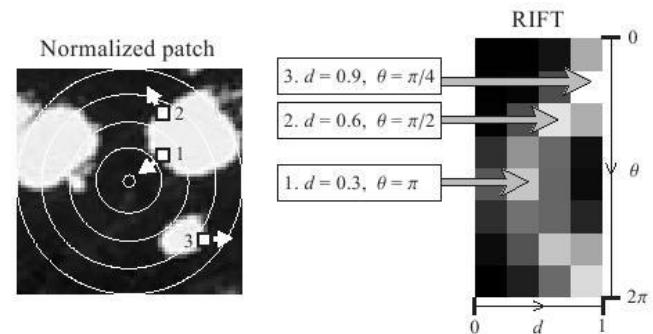
HoG



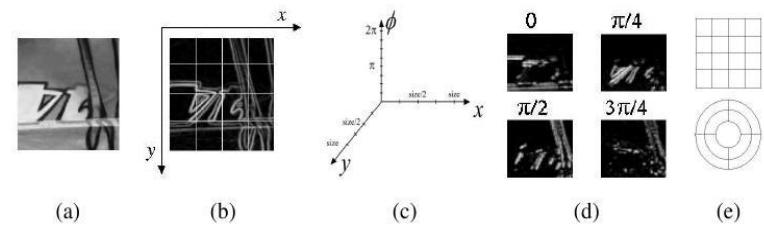
Texton



Spin image



RIFT



GLOH

Learning feature hierarchies with Deep Learning

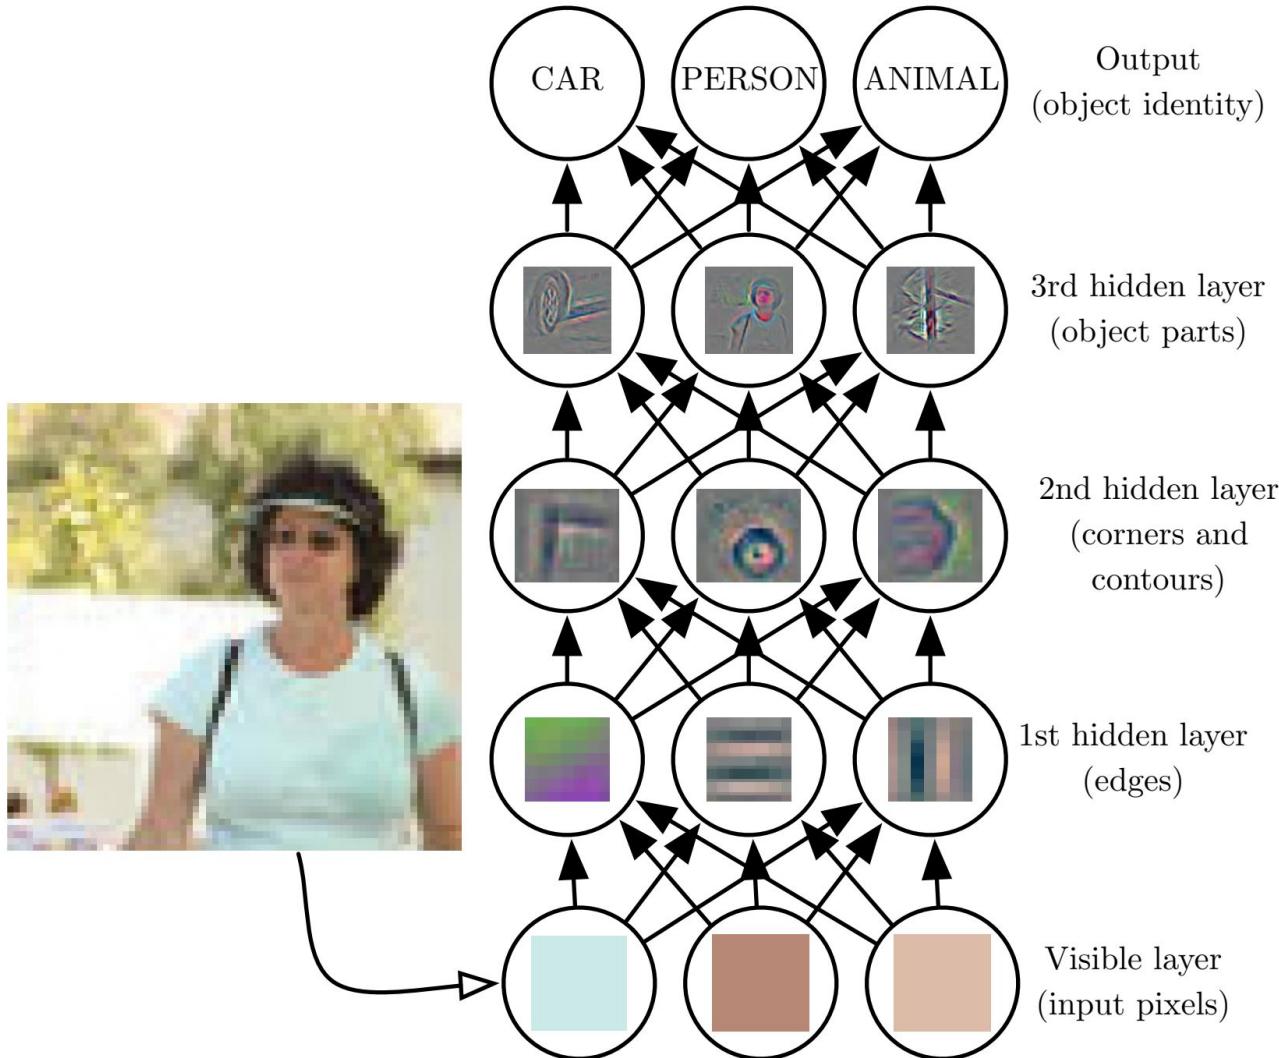
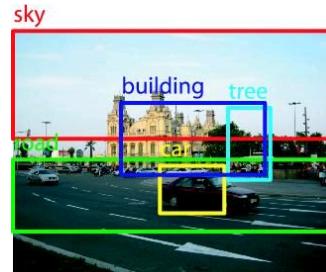


Figure source: Ian Goodfellow, Yoshua Bengio, Aaron Courville. Deep Learning, 2016

ML applications

Examples of ML applications

- Computer vision
- Speech recognition
- Robotics
- Text classification
- Medical image recognition
- Time series prediction/classification
-



WIKIPEDIA
REUTERS 

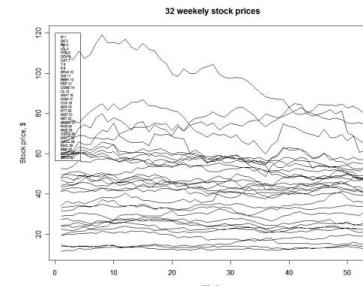


Image classification

- Goal: predict the class of input image

Training dataset
(image, label)



Input

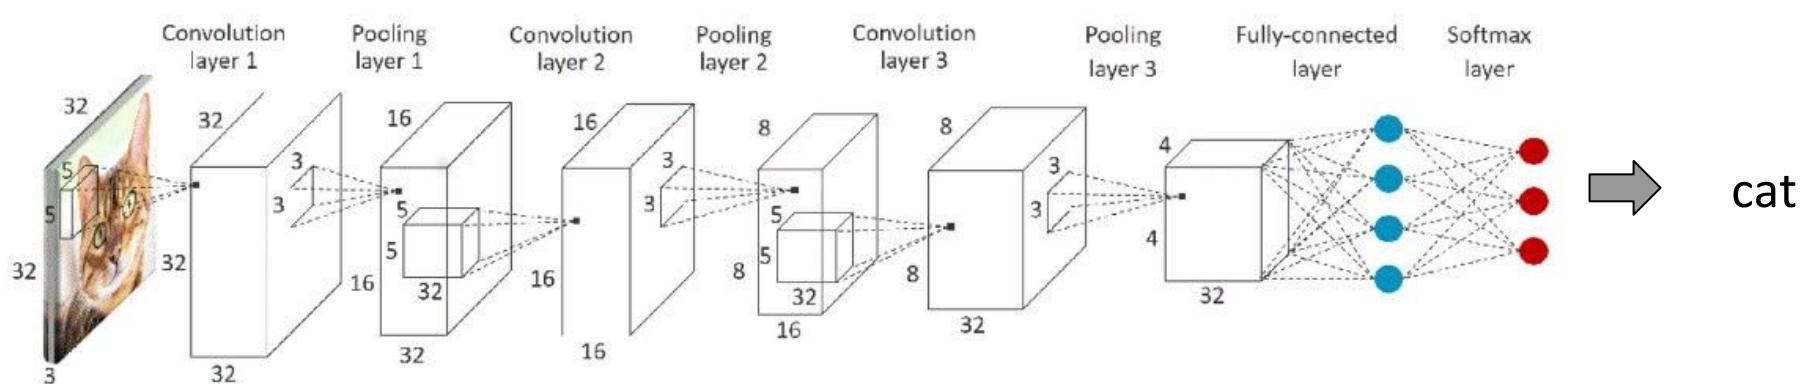
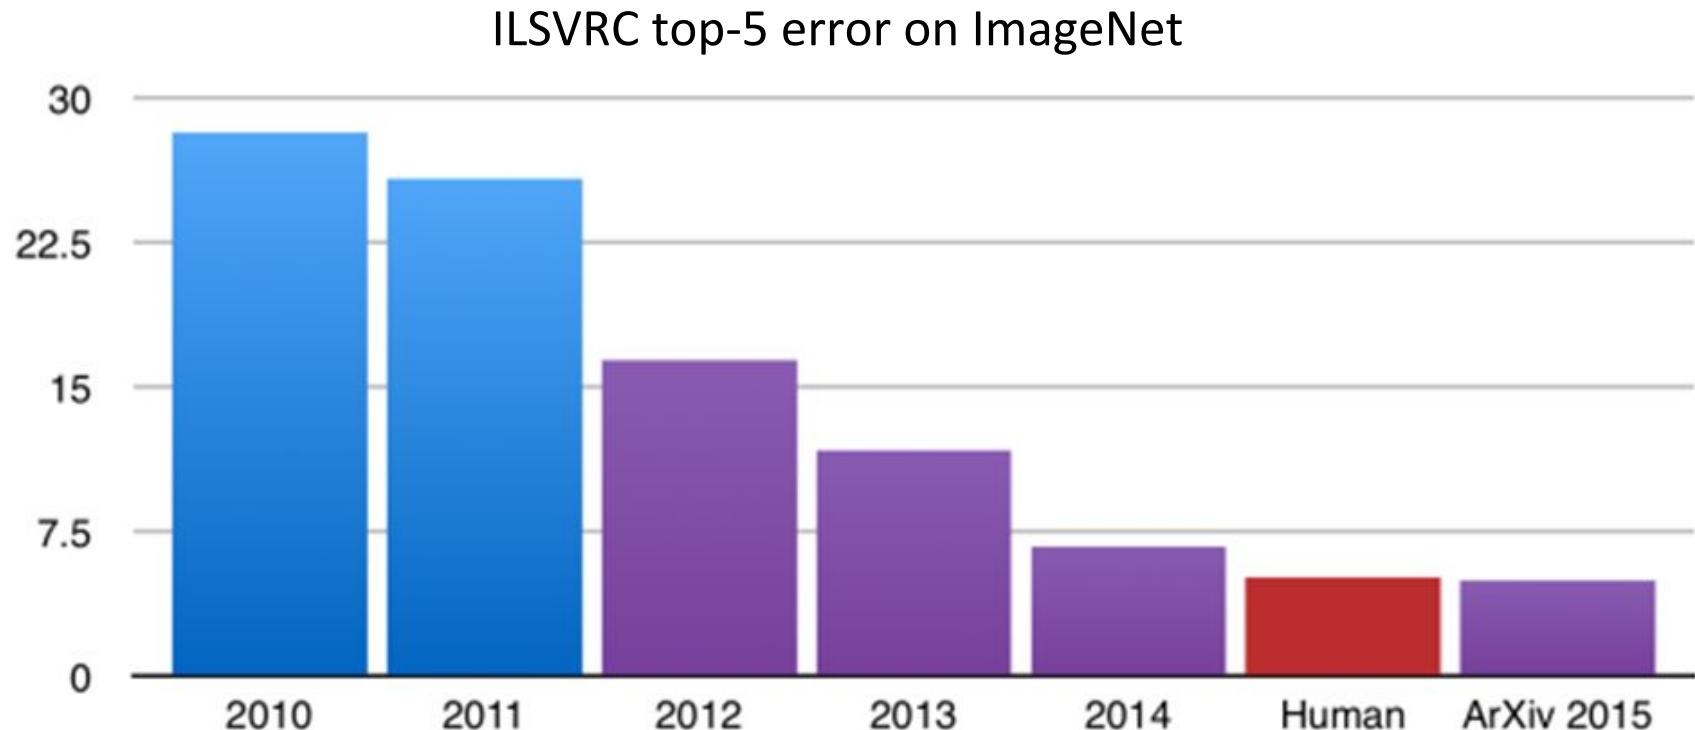


Image classification



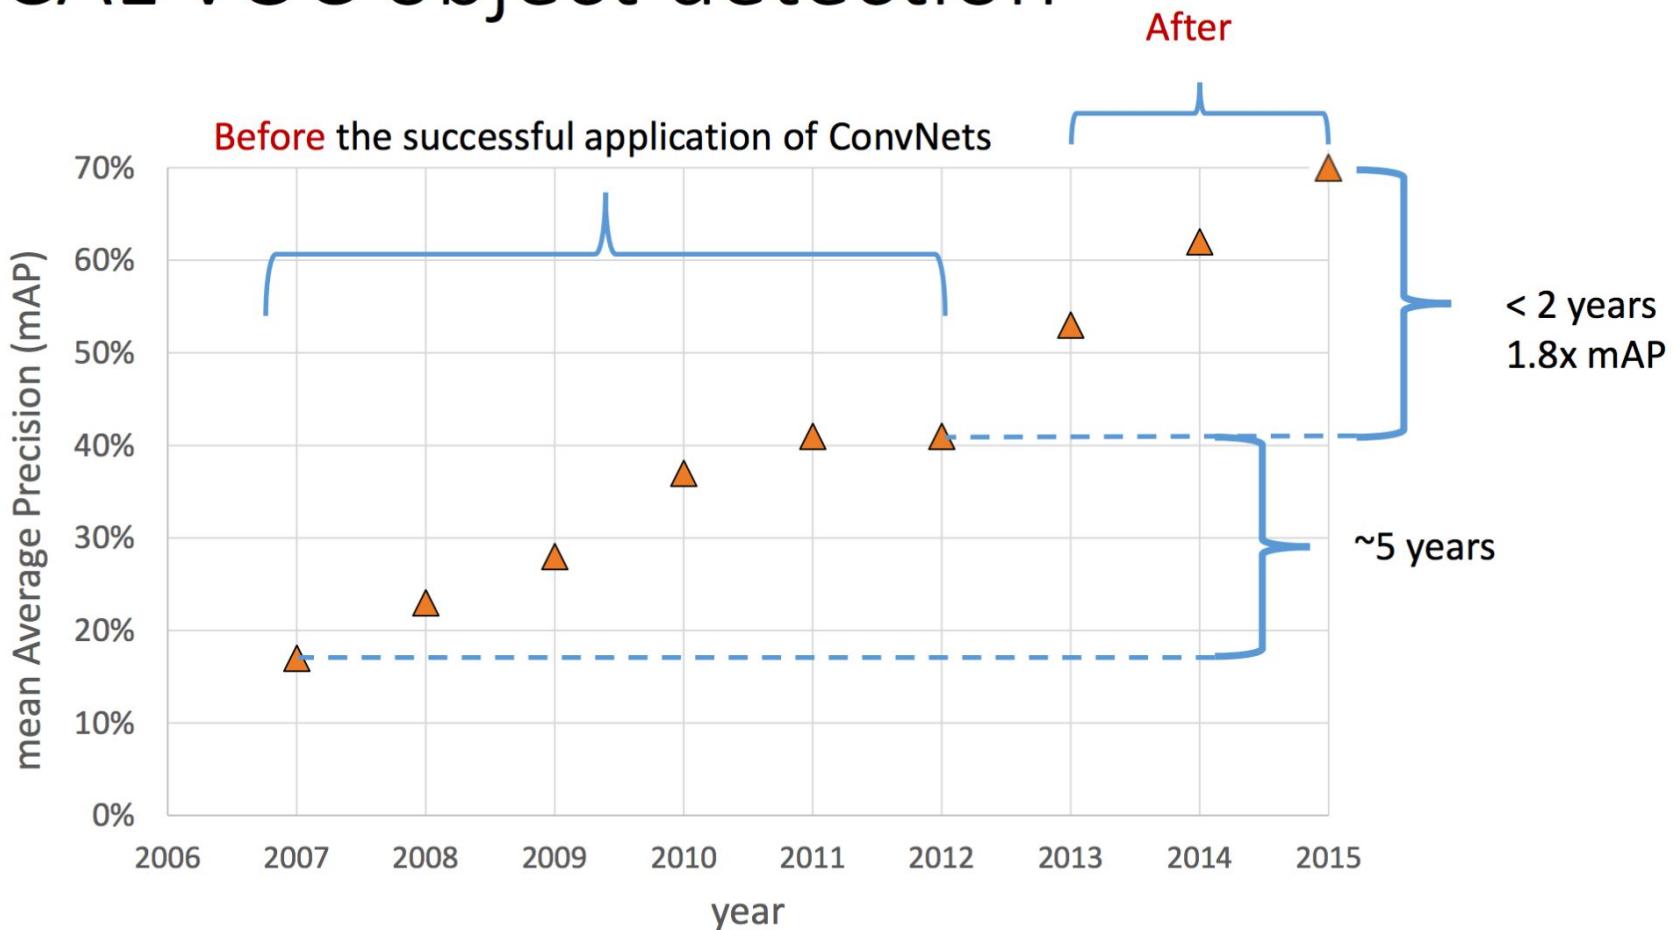
Slide credit: Rob Fergus

Object Detection



Huang et al., Speed/accuracy trade-offs for modern convolutional object detectors. CVPR 2017
figure/code at: https://github.com/tensorflow/models/tree/master/research/object_detection

PASCAL VOC object detection

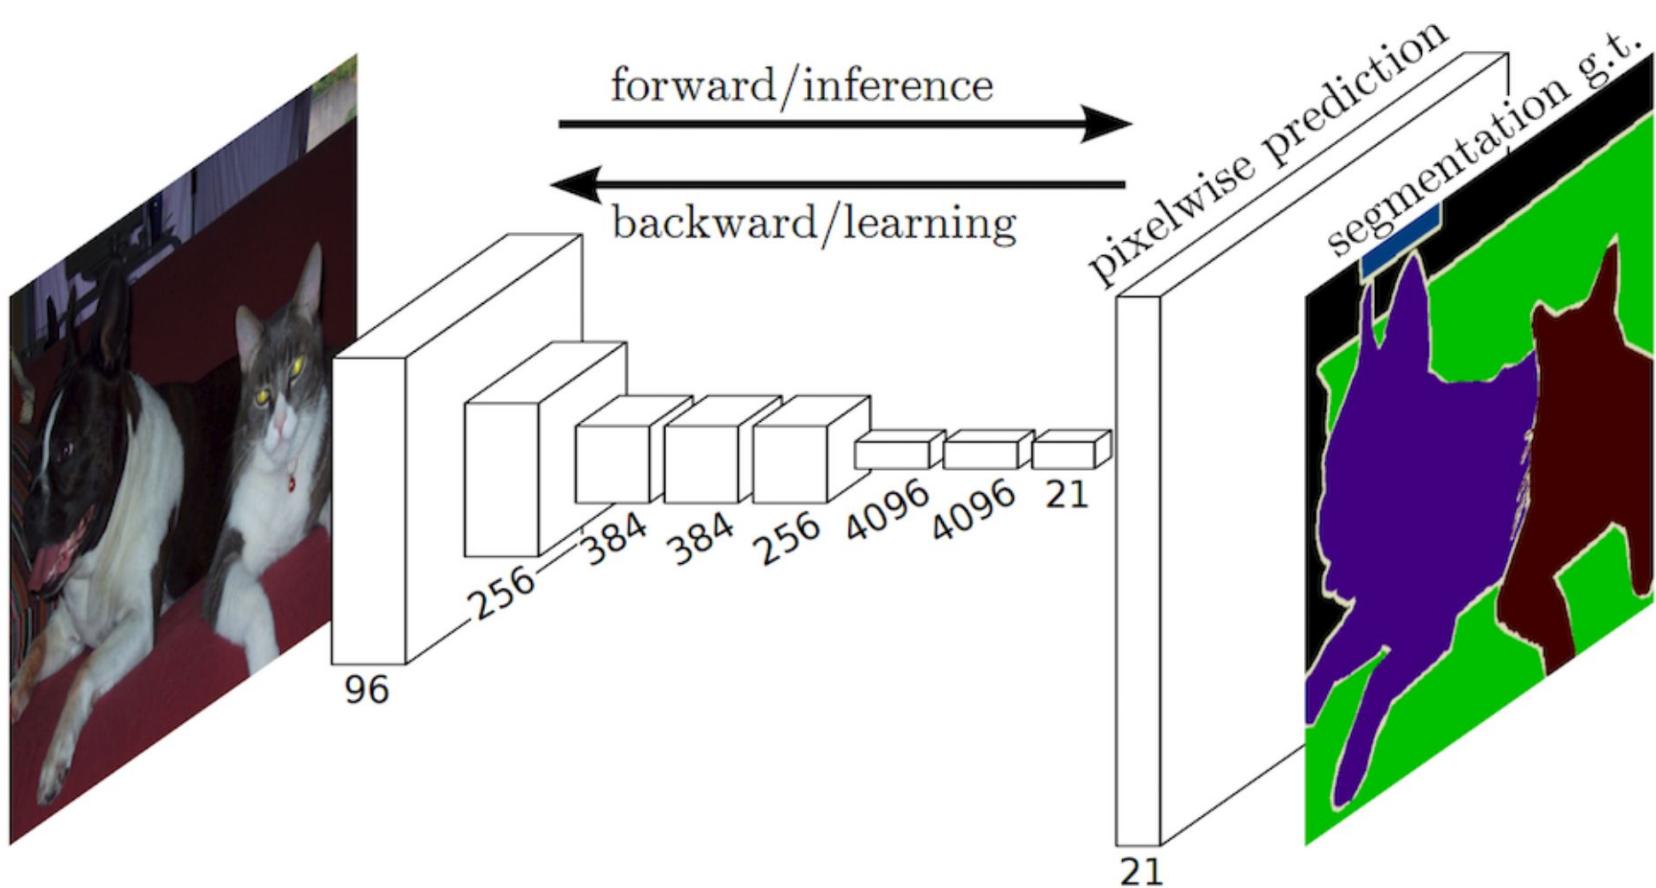


Precision: higher is better

(Figure from Ross Girshick)

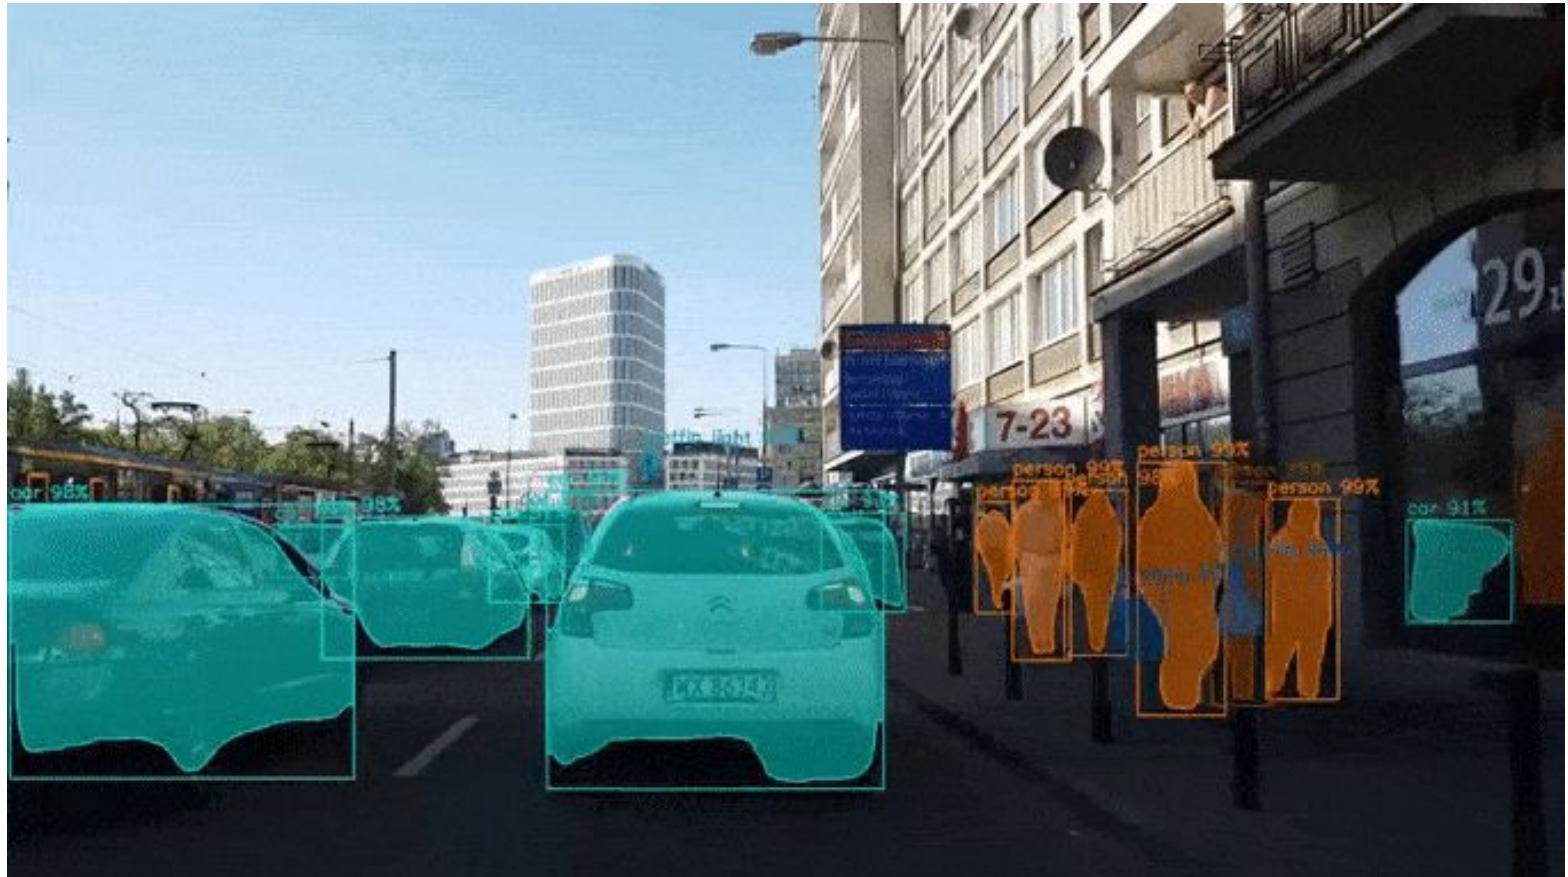
Semantic Segmentation

- Goal: Segment object regions and predict class labels for each region
- Can be formulated as pixel-wise classification



Instance Segmentation

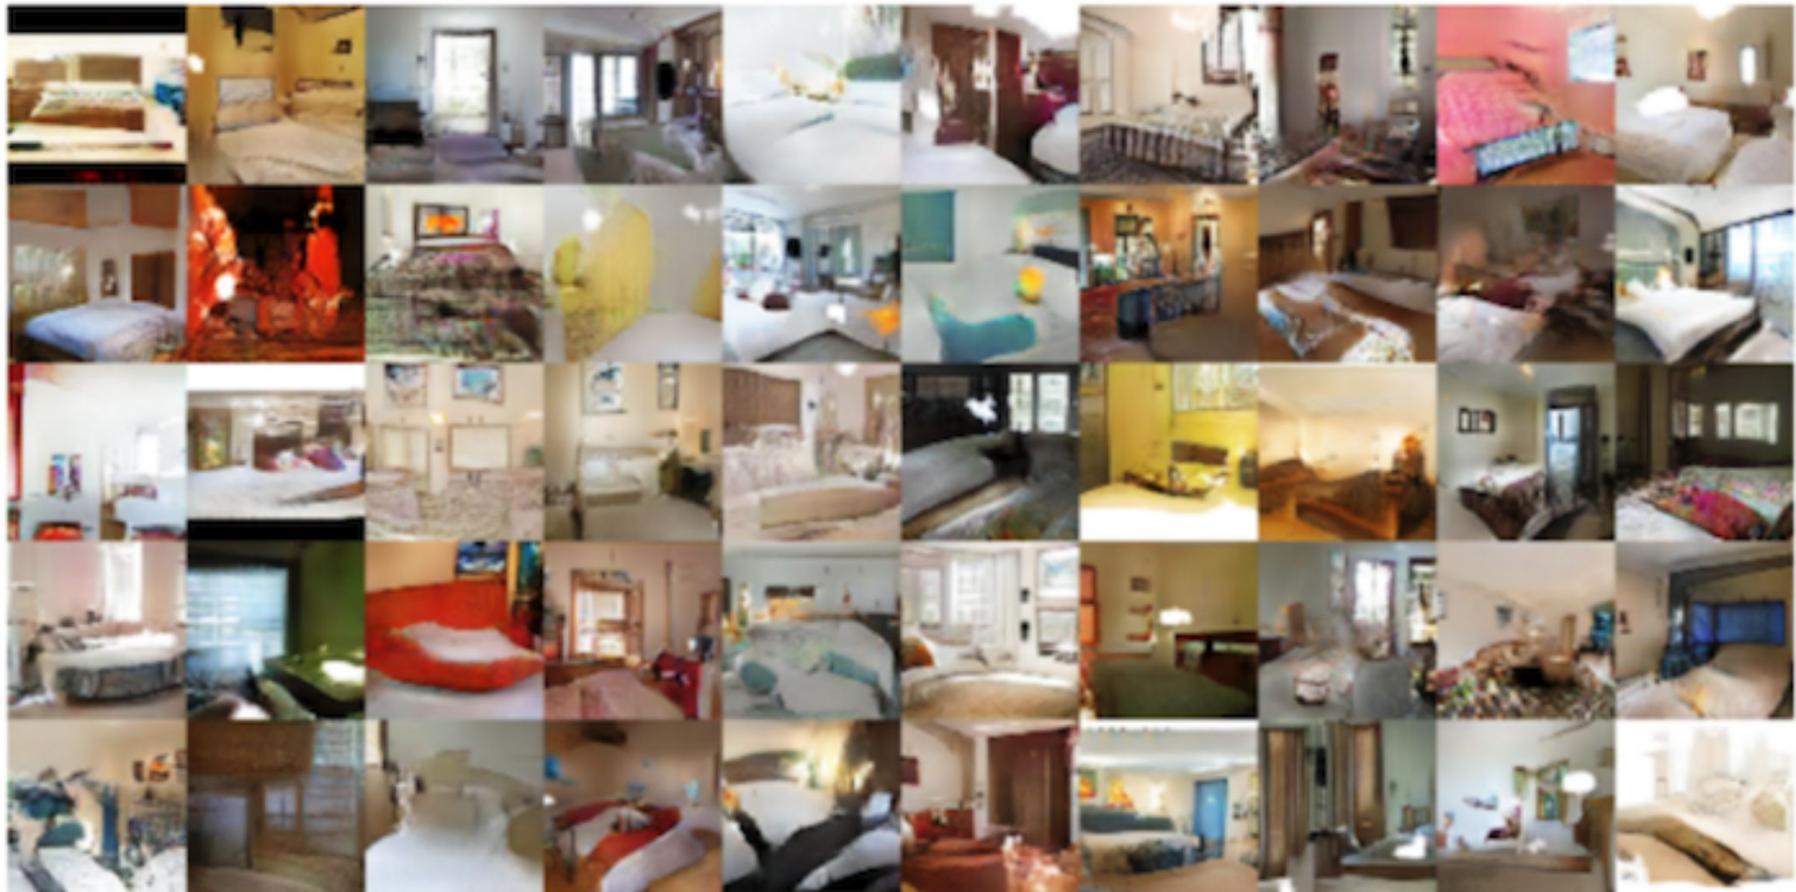
Goal: Perform segmentation (pixel-level masking) for each individual object



Mask R-CNN (He et al., 2017)

Image from: https://github.com/matterport/Mask_RCNN

Image Generation: Generative Adversarial Networks



(Radford et al, "Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks", ICLR, 2016.)

Image Generation: Generative Adversarial Networks



[Karras et al. 2017] Progressive Growing of GANs for Improved Quality, Stability, and Variation

Unsupervised Image-to-Image generation



[CycleGAN: Zhu, Park, Isola & Efros, 2017]

Image Caption Generation

A group of young people playing a game of frisbee.



A herd of elephants walking across a dry grass field.



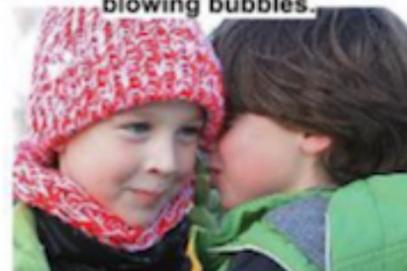
Two hockey players are fighting over the puck.



A close up of a cat laying on a couch.



A little girl in a pink hat is blowing bubbles.



A red motorcycle parked on the side of the road.



A refrigerator filled with lots of food and drinks.



A yellow school bus parked in a parking lot.



Describes without errors

Describes with minor errors

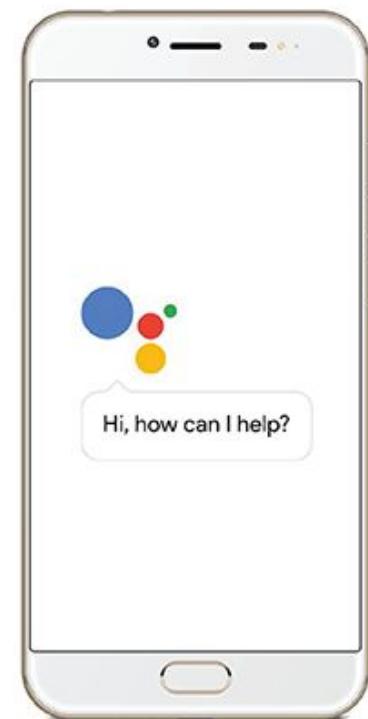
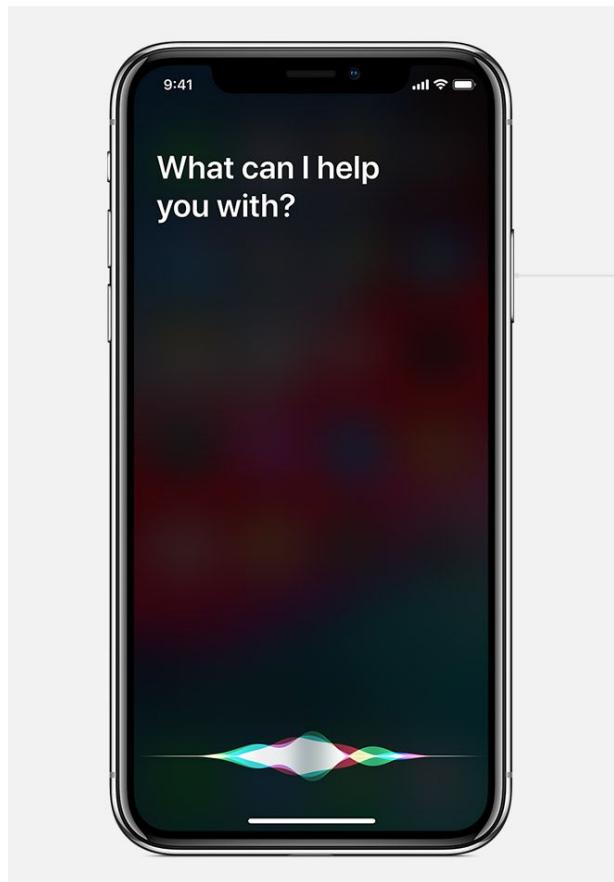
Somewhat related to the image

Unrelated to the image

(Vinyals et al, "Show and Tell: A Neural Image Caption Generator", CVPR, 2015.)

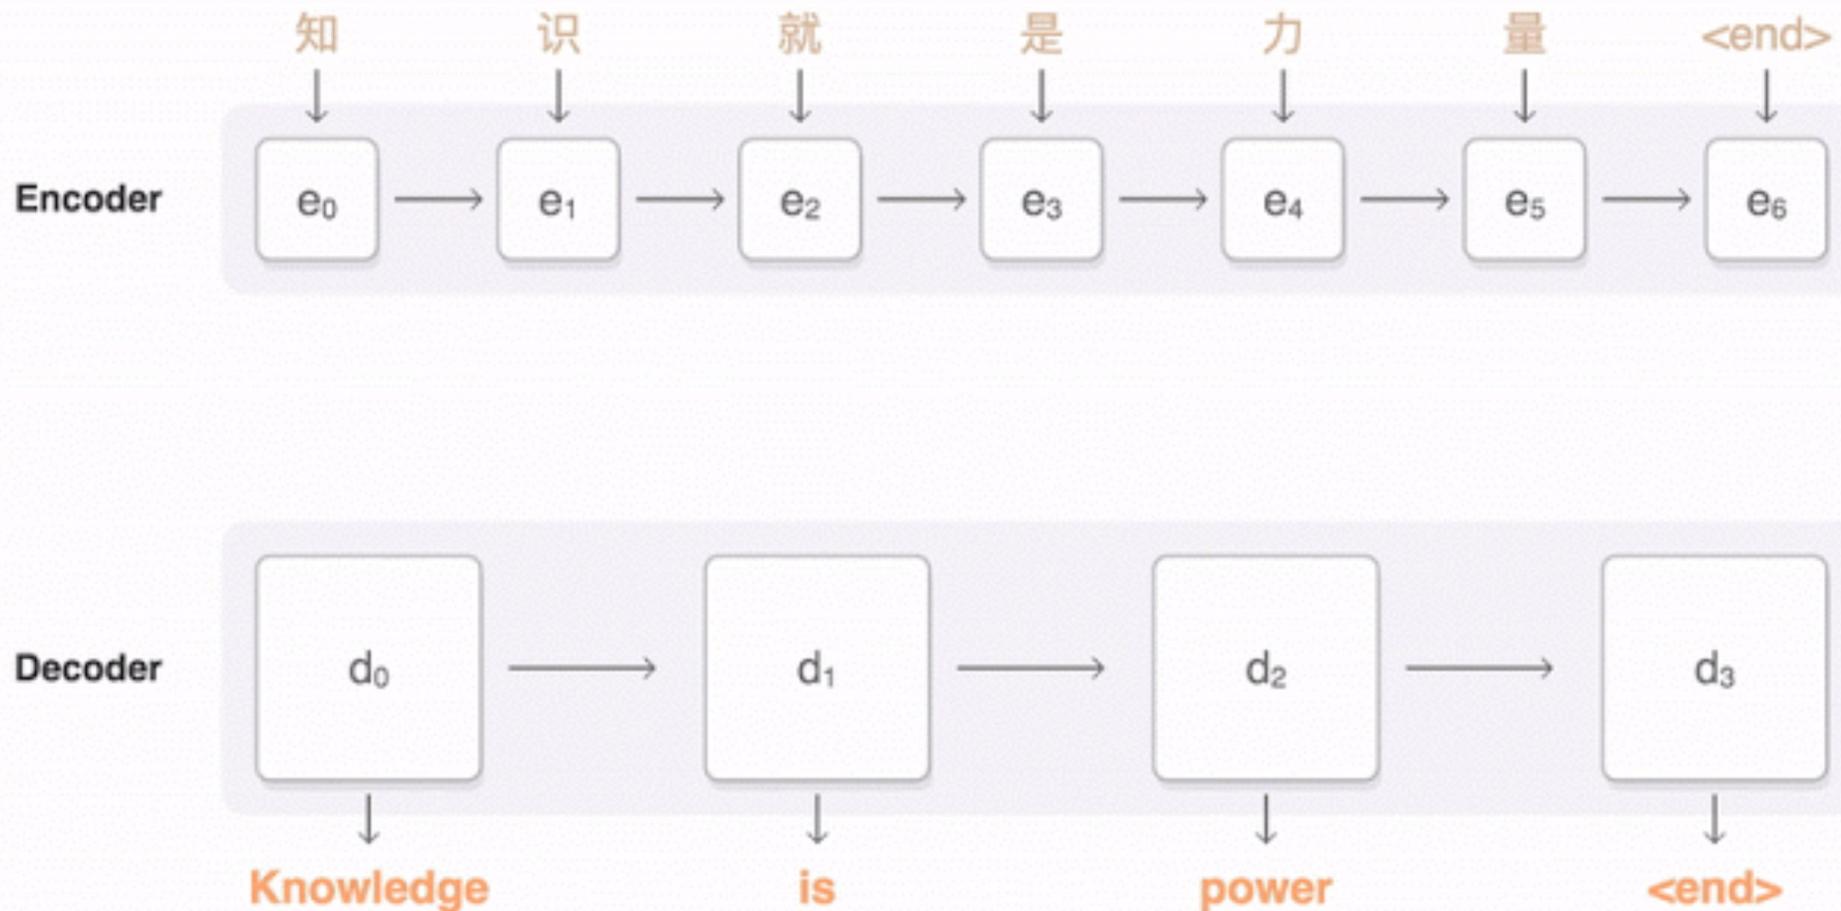
Speech recognition

Siri, Google home, and Google assistant achieves commercial-level performance

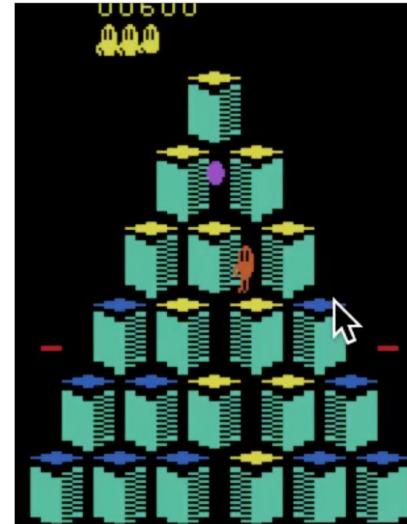
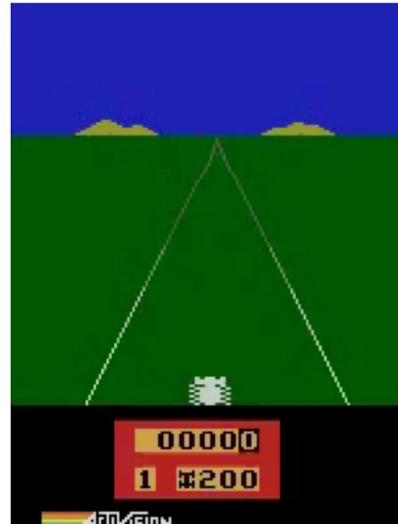


Machine Translation

Google Neural Machine Translation (in production)



RL success stories: playing ATARI games

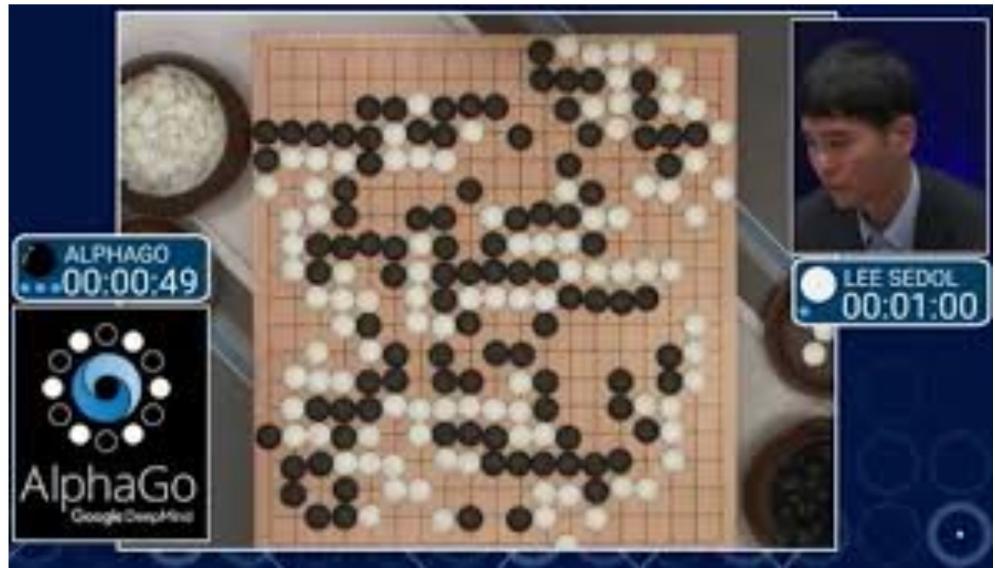


DQN Mnih et al, NIPS 2013 / Nature 2015;

MCTS Guo et al, NIPS 2014; TRPO Schulman, Levine, Moritz, Jordan, Abbeel, ICML 2015;
A3C Mnih et al, ICML 2016; Dueling DQN Wang et al ICML 2016; Double DQN van Hasselt et al, AAAI 2016; Prioritized Experience Replay Schaul et al, ICLR 2016; Bootstrapped DQN Osband et al, 2016; Q-Ensembles Chen et al, 2017; Rainbow Hessel et al, 2017; ...

AlphaGo

- Another breakthrough from Google DeepMind
- Combines Monte-Carlo Tree Search (MCTS) with deep neural networks



AlphaGo Silver et al, Nature 2015

AlphaGoZero Silver et al, Nature 2017

AlphaZero Silver et al, 2017

Tian et al, 2016; Maddison et al, 2014; Clark et al, 2015

OpenAI's 1v1 Dota [2017] and 5v5 [2018]

Super-human agent on a competitive game, enabled by

- Reinforcement learning
- Self-play
- Enough computation

Cooperation emerges



Robot learning



Levine et al., Learning Hand-Eye Coordination for Robotic Grasping. 2016
video: https://www.youtube.com/watch?v=cXaic_k80uM

slide credit: Pieter Abbeel

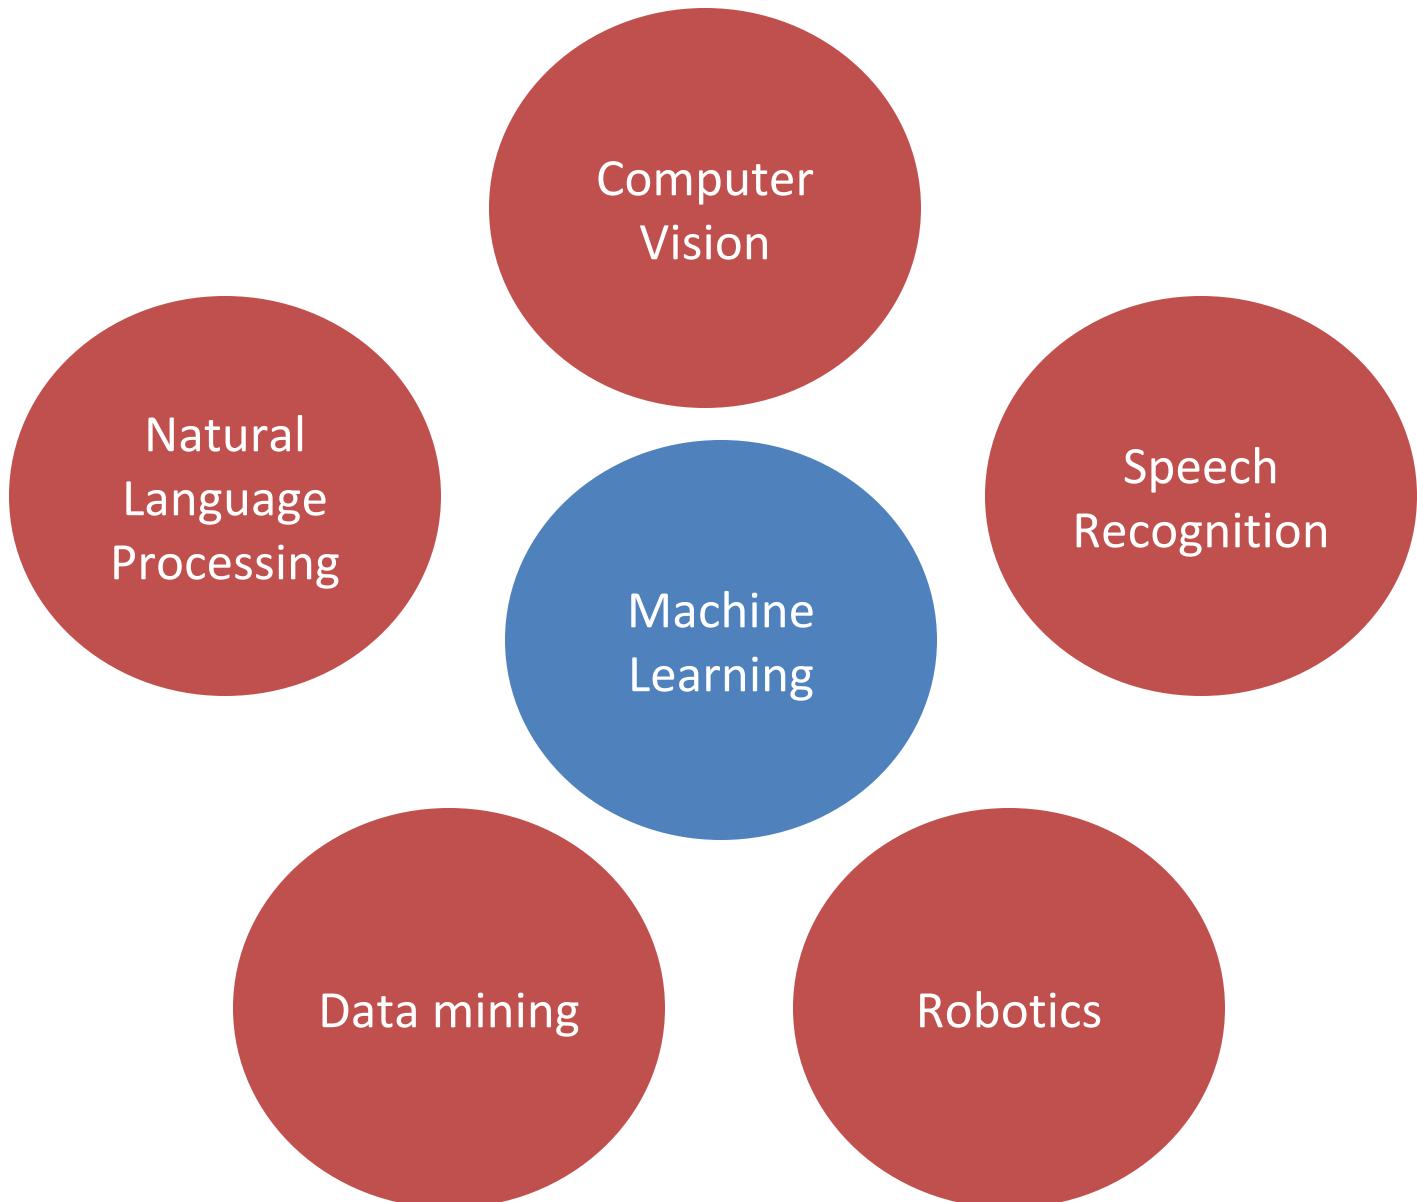
Self-driving cars



<https://youtu.be/O6DRfAC1JXA>

See also: Chris Urmson: How a driverless car sees the road <https://youtu.be/tiwVMrTLUWg>

Machine Learning and other fields



Next class

- Supervised learning
 - Linear regression

Reminder

- Check syllabus at Canvas
- For all questions, please use Piazza (linked to Canvas)

Questions?