

EECS 545: Machine Learning

Lecture 6. Classification 3

Honglak Lee

1/31/2020



Announcements

- HW1 due by today, 23:55
- HW2 due by Feb/11, 23:55
- Google [calendar](#) for the class schedule
 - Subscribe it!

Outline

- Probabilistic generative models
 - Gaussian discriminant analysis (already covered)
 - Naive Bayes
- Discriminant functions
 - Fisher's linear discriminant
 - Perceptron learning algorithm

Recap: Learning the Classifier

- Goal: Learn the distributions $p(C_k | \mathbf{x})$.

(a) Discriminative models: Directly model $p(C_k | \mathbf{x})$ and learn parameters from the training set.

- Logistic regression
- Softmax regression

(b) Generative models: Learn class densities $p(\mathbf{x} | C_k)$ and priors $p(C_k)$

- Gaussian Discriminant Analysis
- Naive Bayes

Recap: Learning the Classifier

- Goal: Learn the distributions $p(C_k | \mathbf{x})$.

(a) Discriminative models: Directly model $p(C_k | \mathbf{x})$ and learn parameters from the training set.

- Logistic regression
- Softmax regression

(b) Generative models: Learn class densities $p(\mathbf{x} | C_k)$ and priors $p(C_k)$

- Gaussian Discriminant Analysis
- **Naive Bayes**

Naive Bayes classifier

Naive Bayes classifier

- Prior distribution:
 - $p(C_k)$: Constant (e.g., Bernoulli)
- Likelihood
 - Naive Bayes assumption: $P(\mathbf{x} | C_k)$ is factorized
(Each coordinate of \mathbf{x} is conditionally independent of other coordinates given the class label)

$$P(x_1, \dots, x_M | C_k) = P(x_1 | C_k) \cdots P(x_M | C_k) = \prod_{j=1}^M P(x_j | C_k)$$

- Classification: use Bayes rule

$$\text{(binary)} \quad P(C_1 | \mathbf{x}) = \frac{P(C_1, \mathbf{x})}{P(\mathbf{x})} = \frac{P(C_1, \mathbf{x})}{P(C_1, \mathbf{x}) + P(C_2, \mathbf{x})}$$

Naive Bayes classifier

- When classifying, we can simply take the MAP (Maximum a Posteriori) estimation:

$$\arg \max_k P(C_k | \mathbf{x}) = \arg \max_k P(C_k, \mathbf{x})$$

Naive Bayes classifier

- When classifying, we can simply take the MAP (Maximum a Posteriori) estimation:

$$\begin{aligned}\arg \max_k P(C_k | \mathbf{x}) &= \arg \max_k P(C_k, \mathbf{x}) \\ &= \arg \max_k P(C_k) P(\mathbf{x} | C_k)\end{aligned}$$

Naive Bayes classifier

- When classifying, we can simply take the MAP (Maximum a Posteriori) estimation:

$$\arg \max_k P(C_k | \mathbf{x}) = \arg \max_k P(C_k, \mathbf{x})$$

$$= \arg \max_k P(C_k) P(\mathbf{x} | C_k)$$

Naive Bayes
assumption



$$= \arg \max_k P(C_k) \prod_{j=1}^M P(x_j | C_k)$$

Example: Spam mail classification

- Label: $y=1$ (spam), $y=0$ (non-spam)
- Features:
 - x_j : j -th word in the mail, where M is the vocabulary size.
 - Multinomial variable (M -dimensional binary vector with only one coordinate with 1)
- Naive Bayes Assumption:
 - Given a class label y , each word in a mail is a independent multinomial variable.

Naive Bayes Spam classifier

- Model

$$P(\text{spam}) = \text{Bernoulli}(\phi)$$

$$P(\text{word}|\text{spam}) = \text{Multinomial}(\mu_1^s, \dots, \mu_M^s)$$

$$P(\text{word}|\text{nospam}) = \text{Multinomial}(\mu_1^{ns}, \dots, \mu_M^{ns})$$

- Goal

: Find ϕ , μ^s , μ^{ns} that best fits the data $\{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(N)}, y^{(N)})\}$

Naive Bayes Spam classifier

- Model

$$P(\text{spam}) = \text{Bernoulli}(\phi)$$

$$P(\text{word}|\text{spam}) = \text{Multinomial}(\mu_1^s, \dots, \mu_M^s)$$

$$P(\text{word}|\text{nospam}) = \text{Multinomial}(\mu_1^{ns}, \dots, \mu_M^{ns})$$

- Goal

: Find ϕ , μ^s , μ^{ns} that best fits the data $\{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(N)}, y^{(N)})\}$

- Likelihood

$$\prod_{i=1}^N P(\mathbf{x}^{(i)}, y^{(i)})$$

Naive Bayes Spam classifier

- Model

$$P(\text{spam}) = \text{Bernoulli}(\phi)$$

$$P(\text{word}|\text{spam}) = \text{Multinomial}(\mu_1^s, \dots, \mu_M^s)$$

$$P(\text{word}|\text{nospam}) = \text{Multinomial}(\mu_1^{ns}, \dots, \mu_M^{ns})$$

- Goal

: Find ϕ , μ^s , μ^{ns} that best fits the data $\{(\mathbf{x}^{(1)}, y^{(1)}), \dots, (\mathbf{x}^{(N)}, y^{(N)})\}$

- Likelihood

$$\begin{aligned} & \prod_{i=1}^N P(\mathbf{x}^{(i)}, y^{(i)}) \\ &= \prod_{i=1}^N P(\mathbf{x}^{(i)} | y^{(i)}) P(y^{(i)}) \\ &= \underbrace{\left(\prod_{i: y^{(i)}=1} P(\mathbf{x}^{(i)} | y^{(i)}) P(y^{(i)}) \right)}_{\text{Spam}} \underbrace{\left(\prod_{i: y^{(i)}=0} P(\mathbf{x}^{(i)} | y^{(i)}) P(y^{(i)}) \right)}_{\text{Non-spam}} \end{aligned}$$

Naive Bayes Spam classifier

- Likelihood - spam

$$\left(\prod_{i:y^{(i)}=1} \underbrace{P(\mathbf{x}^{(i)}|y^{(i)})}_{\text{red}} \underbrace{P(y^{(i)})}_{\text{blue}} \right)$$

$x_k^{(i)}$ \nearrow i-th mail
 \searrow k-th word

- Naive Bayes assumption:

$$P(\text{spam}) = \text{Bernoulli}(\phi)$$

$$P(\text{word}|\text{spam}) = \text{Multinomial}(\mu_1^s, \dots, \mu_M^s)$$

$$\begin{aligned} \underbrace{P(x^{(i)}|y^{(i)} = 1)}_{\text{red}} &= \prod_{k=1}^{\text{len}(x^{(i)})} P(x_k^{(i)}|y^{(i)} = 1) \\ &= \prod_{k=1}^{\text{len}(x^{(i)})} \prod_{j=1}^M (\mu_j^s)^{I(x_k^{(i)} = \text{"}j\text{"th word})} \end{aligned}$$

$$\underbrace{P(y^{(i)} = 1)}_{\text{blue}} = \phi$$

Naive Bayes Spam classifier

- Likelihood - spam (cont')

$$\left(\prod_{i:y^{(i)}=1} P(\mathbf{x}^{(i)}|y^{(i)})P(y^{(i)}) \right)$$
$$= \left(\prod_{i:y^{(i)}=1}^N \prod_{k=1}^{\text{len}(\mathbf{x}^{(i)})} \prod_{j=1}^M (\mu_j^s)^{I(x_k^{(i)} = \text{"}j\text{"th word})} \phi \right)$$

Naive Bayes Spam classifier

- Likelihood - spam (cont')

$$\begin{aligned} & \left(\prod_{i:y^{(i)}=1} P(\mathbf{x}^{(i)}|y^{(i)})P(y^{(i)}) \right) \\ &= \left(\prod_{i:y^{(i)}=1} \prod_{k=1}^{len(x^{(i)})} \prod_{j=1}^M (\mu_j^s)^{I(x_k^{(i)} = \text{"}j\text{"th word})} \phi \right) \\ &= \left(\prod_{i:y^{(i)}=1} \prod_{k=1}^{len(x^{(i)})} \prod_{j=1}^M (\mu_j^s)^{I(x_k^{(i)} = \text{"}j\text{"th word})} \right) \left(\prod_{i:y^{(i)}=1} \phi \right) \end{aligned}$$

Naive Bayes Spam classifier

- Likelihood - spam (cont')

$$\begin{aligned} & \left(\prod_{i:y^{(i)}=1} P(\mathbf{x}^{(i)}|y^{(i)})P(y^{(i)}) \right) \\ &= \left(\prod_{i:y^{(i)}=1} \prod_{k=1}^{\text{len}(\mathbf{x}^{(i)})} \prod_{j=1}^M (\mu_j^s)^{I(x_k^{(i)} = \text{"}j\text{"th word})} \phi \right) \\ &= \left(\prod_{i:y^{(i)}=1} \prod_{k=1}^{\text{len}(\mathbf{x}^{(i)})} \prod_{j=1}^M (\mu_j^s)^{I(x_k^{(i)} = \text{"}j\text{"th word})} \right) \left(\prod_{i:y^{(i)}=1} \phi \right) \\ &= \left(\prod_{j=1}^M (\mu_j^s)^{\sum_{i:y^{(i)}=1} \sum_{k=1}^{\text{len}(\mathbf{x}^{(i)})} I(x_k^{(i)} = \text{"}j\text{"th word})} \right) \left(\prod_{i:y^{(i)}=1} \phi \right) \end{aligned}$$

Naive Bayes Spam classifier

- Likelihood - spam (cont')

$$\begin{aligned}
 & \left(\prod_{i:y^{(i)}=1} P(\mathbf{x}^{(i)}|y^{(i)})P(y^{(i)}) \right) \\
 &= \left(\prod_{i:y^{(i)}=1} \prod_{k=1}^{\text{len}(\mathbf{x}^{(i)})} \prod_{j=1}^M (\mu_j^s)^{I(x_k^{(i)} = \text{"}j\text{"th word})} \phi \right) \\
 &= \left(\prod_{i:y^{(i)}=1} \prod_{k=1}^{\text{len}(\mathbf{x}^{(i)})} \prod_{j=1}^M (\mu_j^s)^{I(x_k^{(i)} = \text{"}j\text{"th word})} \right) \left(\prod_{i:y^{(i)}=1} \phi \right) \\
 &= \left(\prod_{j=1}^M (\mu_j^s)^{\sum_{i:y^{(i)}=1} \sum_{k=1}^{\text{len}(\mathbf{x}^{(i)})} I(x_k^{(i)} = \text{"}j\text{"th word})} \right) \left(\prod_{i:y^{(i)}=1} \phi \right) \\
 &= \left(\prod_{j=1}^M (\mu_j^s)^{N_j^{\text{spam}}} \right) \phi^{N^{\text{spam}}}
 \end{aligned}$$

Maximum likelihood estimation

- Putting together:

$$\prod_{i=1}^N P(\mathbf{x}^{(i)}, y^{(i)})$$
$$= \left(\prod_{i: y^{(i)}=1} P(\mathbf{x}^{(i)} | y^{(i)}) P(y^{(i)}) \right) \left(\prod_{i: y^{(i)}=0} P(\mathbf{x}^{(i)} | y^{(i)}) P(y^{(i)}) \right)$$

Maximum likelihood estimation

- Putting together:

$$\begin{aligned} & \prod_{i=1}^N P(\mathbf{x}^{(i)}, y^{(i)}) \\ &= \left(\prod_{i: y^{(i)}=1} P(\mathbf{x}^{(i)} | y^{(i)}) P(y^{(i)}) \right) \left(\prod_{i: y^{(i)}=0} P(\mathbf{x}^{(i)} | y^{(i)}) P(y^{(i)}) \right) \\ &= \left(\phi^{N^{spam}} \prod_{word\ j} (\mu_j^s)^{N_j^{spam}} \right) \left((1 - \phi)^{N^{nonsпам}} \prod_{word\ j} (\mu_j^{ns})^{N_j^{nonsпам}} \right) \end{aligned}$$

Maximum likelihood estimation

- Putting together:

$$\begin{aligned} & \prod_{i=1}^N P(\mathbf{x}^{(i)}, y^{(i)}) \\ &= \left(\prod_{i: y^{(i)}=1} P(\mathbf{x}^{(i)} | y^{(i)}) P(y^{(i)}) \right) \left(\prod_{i: y^{(i)}=0} P(\mathbf{x}^{(i)} | y^{(i)}) P(y^{(i)}) \right) \\ &= \left(\phi^{N^{spam}} \prod_{word\ j} (\mu_j^s)^{N_j^{spam}} \right) \left((1 - \phi)^{N^{nonspam}} \prod_{word\ j} (\mu_j^{ns})^{N_j^{nonspam}} \right) \end{aligned}$$

- Log-likelihood

$$\begin{aligned} & \log P(\mathcal{D}) \\ &= \log \prod_{i=1}^N P(x^{(i)}, y^{(i)}) \\ &= N^{spam} \log \phi + \sum_{word\ j} N_j^{spam} \log \mu_j^s + N^{nonspam} \log(1 - \phi) + \sum_{word\ j} N_j^{nonspam} \log \mu_j^{ns} \end{aligned}$$

Maximum likelihood estimation

- Log-likelihood

$$\begin{aligned} & \log P(\mathcal{D}) \\ = & \log \prod_{i=1}^N P(x^{(i)}, y^{(i)}) \\ = & N^{spam} \log \phi + \sum_{word\ j} N_j^{spam} \log \mu_j^s + N^{nonspam} \log(1 - \phi) + \sum_{word\ j} N_j^{nonspam} \log \mu_j^{ns} \end{aligned}$$

- Maximum-likelihood

- Take the derivative of log-likelihood w.r.t. the parameters, and set it to zero.

Maximum likelihood estimation

- From $\frac{\partial l}{\partial \phi} = \frac{1}{\phi} N^{spam} - \frac{1}{1-\phi} N^{nonspam} = 0$
 - We get $\phi = \frac{N^{spam}}{N^{spam} + N^{nonspam}}$

- Make the parameters μ independent:

$$\sum_{word\ j=1}^M N_j^{spam} \log \mu_j^s = \sum_{word\ j=1}^{M-1} N_j^{spam} \log \mu_j^s + N_M^{spam} \log(1 - \sum_{j=1}^{M-1} \mu_j^s)$$

$$\frac{\partial}{\partial \mu_j^s} \left(\sum_{word\ j=1}^M N_j^{spam} \log \mu_j^s \right) = \frac{N_j^{spam}}{\mu_j^s} - \frac{N_M^{spam}}{1 - \sum_{j=1}^{M-1} \mu_j^s} = 0$$

Maximum likelihood estimation

- From $\frac{\partial l}{\partial \phi} = \frac{1}{\phi} N^{spam} - \frac{1}{1-\phi} N^{nonspam} = 0$
 - We get $\phi = \frac{N^{spam}}{N^{spam} + N^{nonspam}}$

- Make the parameters μ independent:

$$\sum_{word\ j=1}^M N_j^{spam} \log \mu_j^s = \sum_{word\ j=1}^{M-1} N_j^{spam} \log \mu_j^s + N_M^{spam} \log(1 - \sum_{j=1}^{M-1} \mu_j^s)$$

$$\frac{\partial}{\partial \mu_j^s} \left(\sum_{word\ j=1}^M N_j^{spam} \log \mu_j^s \right) = \frac{N_j^{spam}}{\mu_j^s} - \frac{N_M^{spam}}{1 - \sum_{j=1}^{M-1} \mu_j^s} = 0$$

$$\frac{N_j^{spam}}{\mu_j^s} = constant, \forall j$$

We finally get

$$\mu_j^s = \frac{N_j^{spam}}{\sum_j N_j^{spam}}$$

Maximum likelihood estimation

- Summary:

$$P(spam) = \phi = \frac{N^{spam}}{N^{spam} + N^{nonspam}}$$

$$P(word = j|spam) = \mu_j^s = \frac{N_j^{spam}}{\sum_j N_j^{spam}}$$

$$P(word = j|non - spam) = \mu_j^{ns} = \frac{N_j^{nonspam}}{\sum_j N_j^{nonspam}}$$

Recall:

N^{spam} : total # examples for spam

$N^{nonspam}$: total # examples for non-spam

N_j^{spam} : total # word j from the entire spam emails

$N_j^{nonspam}$: total # word j from the entire nonspam emails

Laplace smoothing

- Maximum likelihood is problematic when a specific word count is 0
 - Leads to probability of 0!
- Solution: Put “imaginary” counts for each word
 - prevent zero probability estimates (overfitting)!
 - E.g.: Adding “1” as imaginary count for each word

$$P(spam) = \phi = \frac{N^{spam}}{N^{spam} + N^{nonspam}}$$

$$P(word = j|spam) = \mu_j^s = \frac{N_j^{spam} + 1}{\sum_j N_j^{spam} + M}$$

$$P(word = j|non - spam) = \mu_j^{ns} = \frac{N_j^{nonspam} + 1}{\sum_j N_j^{nonspam} + M}$$


Discriminant functions

Linear Discriminant functions: Discriminating two classes

- Specify a weight vector \mathbf{w} and a bias w_0

$$h(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$

- Assign \mathbf{x} to C_1 if $h(\mathbf{x}) \geq 0$ and to C_0 otherwise.
- Q: How to pick \mathbf{w} ?



Linear Discriminant functions: Discriminating $K > 2$ classes

- Instead each class C_k gets its own function


$$h_k(\mathbf{x}) = \mathbf{w}_k^T \mathbf{x} + w_{k,0}$$

- Assign \mathbf{x} to C_k if

$$h_k(\mathbf{x}) > h_j(\mathbf{x}) \text{ for all } j \neq k$$

- The decision regions are convex polyhedra.



Decision Regions



- Decision regions are convex, with piecewise linear boundaries.

How do we set the weights w ?

- How about w that minimizes squared error?
 - Label y versus linear prediction $h(w)$.
 - Least squares is too sensitive to outliers. (why?)



Read Bishop book

Learning linear discriminant functions (this lecture)


- Fisher's linear discriminant
- Perceptron learning algorithm

Fisher's Linear Discriminant

- Use w to project x to one dimension.

$$\text{if } w^T x \geq -w_0 \text{ then } C_1 \text{ else } C_0$$



- Select w that best separates the classes.
- By “separating”, the algorithm simultaneously
 - maximizes between-class variances
 - minimizes within-class variances



Read Bishop book

Fisher's Linear Discriminant

- Maximizing separation alone doesn't work.
 - Minimizing class variance is a big help.



See Bishop book

Objective function

- We want to maximize the “distance between classes”

$$\underline{m_2} - m_1 \equiv \mathbf{w}^T (\underline{\mathbf{m}_2} - \mathbf{m}_1) \quad \text{where } \mathbf{m}_k = \frac{1}{N_k} \sum_{n \in C_k} \mathbf{x}_n$$

Projected mean

Mean

See Bishop book

Objective function

- We want to maximize the “distance between classes”

$$\underline{m_2} - m_1 \equiv \mathbf{w}^T (\underline{\mathbf{m}_2} - \mathbf{m}_1) \quad \text{where } \mathbf{m}_k = \frac{1}{N_k} \sum_{n \in C_k} \mathbf{x}_n$$

Projected mean

Mean

- While minimizing the “distance within each class”

$$s_1^2 + s_2^2 \equiv \sum_{n \in C_1} (\mathbf{w}^T \mathbf{x}_n - m_1)^2 + \sum_{n \in C_2} (\mathbf{w}^T \mathbf{x}_n - m_2)^2$$

See Bishop book

Objective function

- We want to maximize the “distance between classes”

$$\underline{m_2} - m_1 \equiv \mathbf{w}^T (\underline{\mathbf{m}_2} - \mathbf{m}_1) \quad \text{where } \mathbf{m}_k = \frac{1}{N_k} \sum_{n \in C_k} \mathbf{x}_n$$

Projected mean

Mean

- While minimizing the “distance within each class”

$$s_1^2 + s_2^2 \equiv \sum_{n \in C_1} (\mathbf{w}^T \mathbf{x}_n - m_1)^2 + \sum_{n \in C_2} (\mathbf{w}^T \mathbf{x}_n - m_2)^2$$

- Objective function: $J(\mathbf{w}) = \frac{(m_2 - m_1)^2}{s_1^2 + s_2^2}$

See Bishop book

Derivation of objective

- Numerator: $m_2 - m_1 = \mathbf{w}^T(\mathbf{m}_2 - \mathbf{m}_1)$:
 - $\|m_2 - m_1\|^2 = \mathbf{w}^T(\mathbf{m}_2 - \mathbf{m}_1)(\mathbf{m}_2 - \mathbf{m}_1)^T \mathbf{w}$

See Bishop book

Derivation of objective

- Numerator: $m_2 - m_1 = \mathbf{w}^T(\mathbf{m}_2 - \mathbf{m}_1)$:
 - $\|m_2 - m_1\|^2 = \mathbf{w}^T(\mathbf{m}_2 - \mathbf{m}_1)(\mathbf{m}_2 - \mathbf{m}_1)^T \mathbf{w}$
- Denominator:
 - $s_k^2 = \sum_{n \in C_k} (\mathbf{w}^T \mathbf{x}_n - m_k)^2$
 $= \sum_{n \in C_k} \mathbf{w}^T (\mathbf{x}_n - \mathbf{m}_k)(\mathbf{x}_n - \mathbf{m}_k)^T \mathbf{w}$
 - $s_1^2 + s_2^2 = \mathbf{w}^T \left[\sum_{k=1,2} \sum_{n \in C_k} (\mathbf{x}_n - \mathbf{m}_k)(\mathbf{x}_n - \mathbf{m}_k)^T \right] \mathbf{w}$

See Bishop book

Derivation of objective

- Numerator: $m_2 - m_1 = \mathbf{w}^T(\mathbf{m}_2 - \mathbf{m}_1)$:
 - $\|m_2 - m_1\|^2 = \mathbf{w}^T(\mathbf{m}_2 - \mathbf{m}_1)(\mathbf{m}_2 - \mathbf{m}_1)^T \mathbf{w}$
- Denominator:
 - $s_k^2 = \sum_{n \in C_k} (\mathbf{w}^T \mathbf{x}_n - m_k)^2$
 $= \sum_{n \in C_k} \mathbf{w}^T (\mathbf{x}_n - \mathbf{m}_k)(\mathbf{x}_n - \mathbf{m}_k)^T \mathbf{w}$
 - $s_1^2 + s_2^2 = \mathbf{w}^T \left[\sum_{k=1,2} \sum_{n \in C_k} (\mathbf{x}_n - \mathbf{m}_k)(\mathbf{x}_n - \mathbf{m}_k)^T \right] \mathbf{w}$
- After definition of terms, we get
$$J(\mathbf{w}) = \frac{\mathbf{w}^T S_B \mathbf{w}}{\mathbf{w}^T S_W \mathbf{w}}$$
 - Solution: $\mathbf{w} \propto S_W^{-1}(\mathbf{m}_2 - \mathbf{m}_1)$

See Bishop book

The Perceptron

- A “generalized linear function”

$$h(\mathbf{x}) = f(\mathbf{w}^T \phi(\mathbf{x}))$$

- Where

$$f(a) = \begin{cases} +1, & a \geq 0 \\ -1, & a < 0 \end{cases}$$

- Uses target code: $y=+1$ for C_1 , $y=-1$ for C_2 .
- Means we always want:

$$\mathbf{w}^T \phi(\mathbf{x}_n) y_n > 0$$

The Perceptron Criterion

- Only count errors from misclassified points:

$$E_P(\mathbf{w}) = - \sum_{\mathbf{x}_n \in \mathcal{M}} \mathbf{w}^T \phi(\mathbf{x}) y_n$$



- where \mathcal{M} are the misclassified points.
- Stochastic gradient descent:
 - Update the weight vector according to the misclassified points:

$$\mathbf{w}^{(\tau+1)} = \mathbf{w}^{(\tau)} - \eta \nabla E_P(\mathbf{w}) = \mathbf{w}^{(\tau)} + \eta \phi(\mathbf{x}) y_n$$

Note: update only for misclassified examples



Perceptron Learning (1)

- If \mathbf{x}_n is misclassified, add $\phi(\mathbf{x}_n)$ into w .



Perceptron Learning (2)

- If \mathbf{x}_n is misclassified, add $\phi(\mathbf{x}_n)$ into w .



Perceptron Learning

- Perceptron Convergence Theorem:
 - If there exists an exact solution (i.e., if the training data is linearly separable)
 - then the learning algorithm will find it in a finite number of steps.
- Limitations of perceptron learning:
 - The convergence can be very slow.
 - If dataset is not linearly separable, it won't converge.
 - Does not generalize well to $K > 2$ classes.

Next class

- Kernel methods